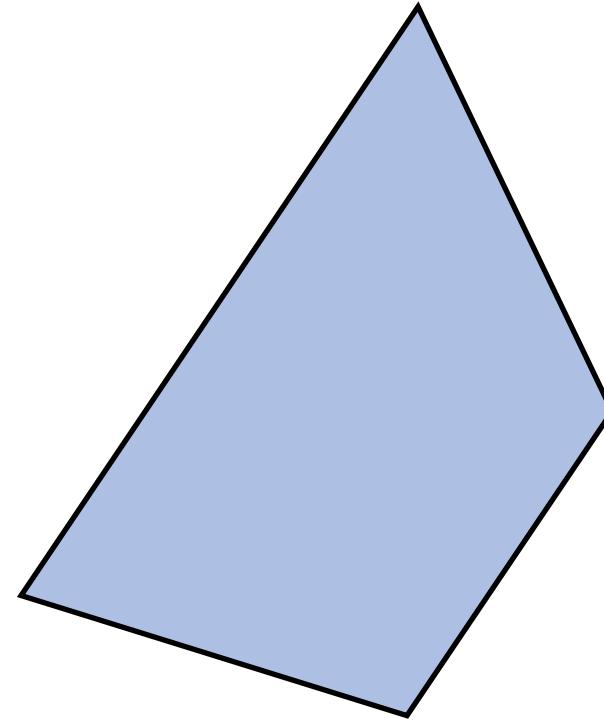
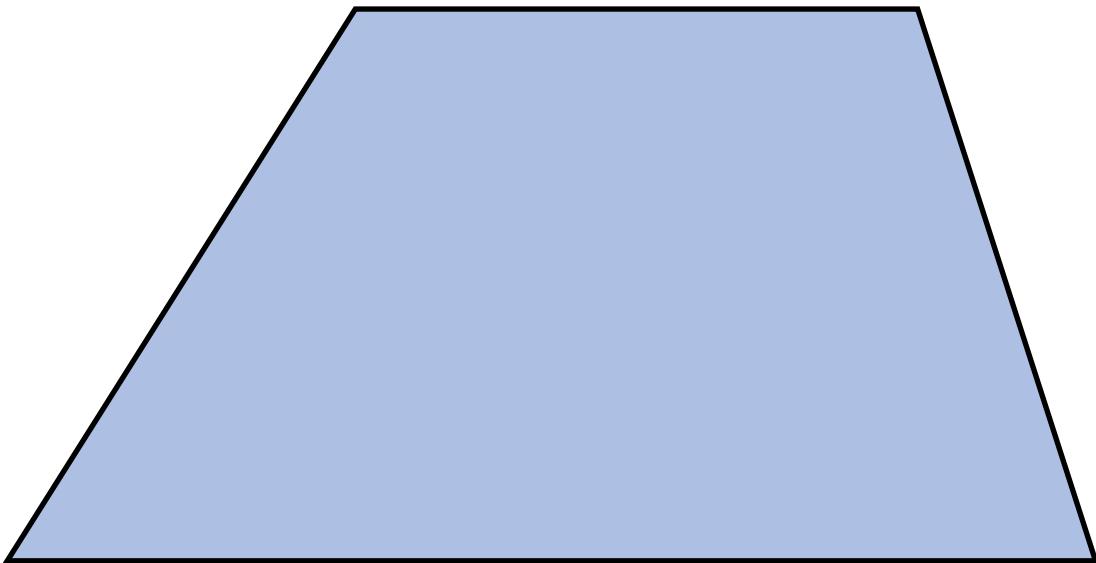


# Corner Detection



Pick a point in the image.  
Find it again in the next image.

*What type of feature would you select?*  
a corner

# Why detect corners?

Image alignment (homography, fundamental matrix)

3D reconstruction

Motion tracking

Object recognition

Indexing and database retrieval

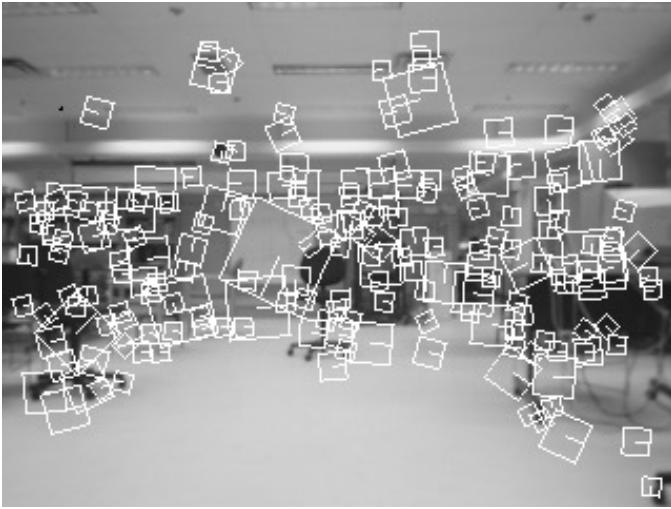
Robot navigation

And so on....

# Location Recognition

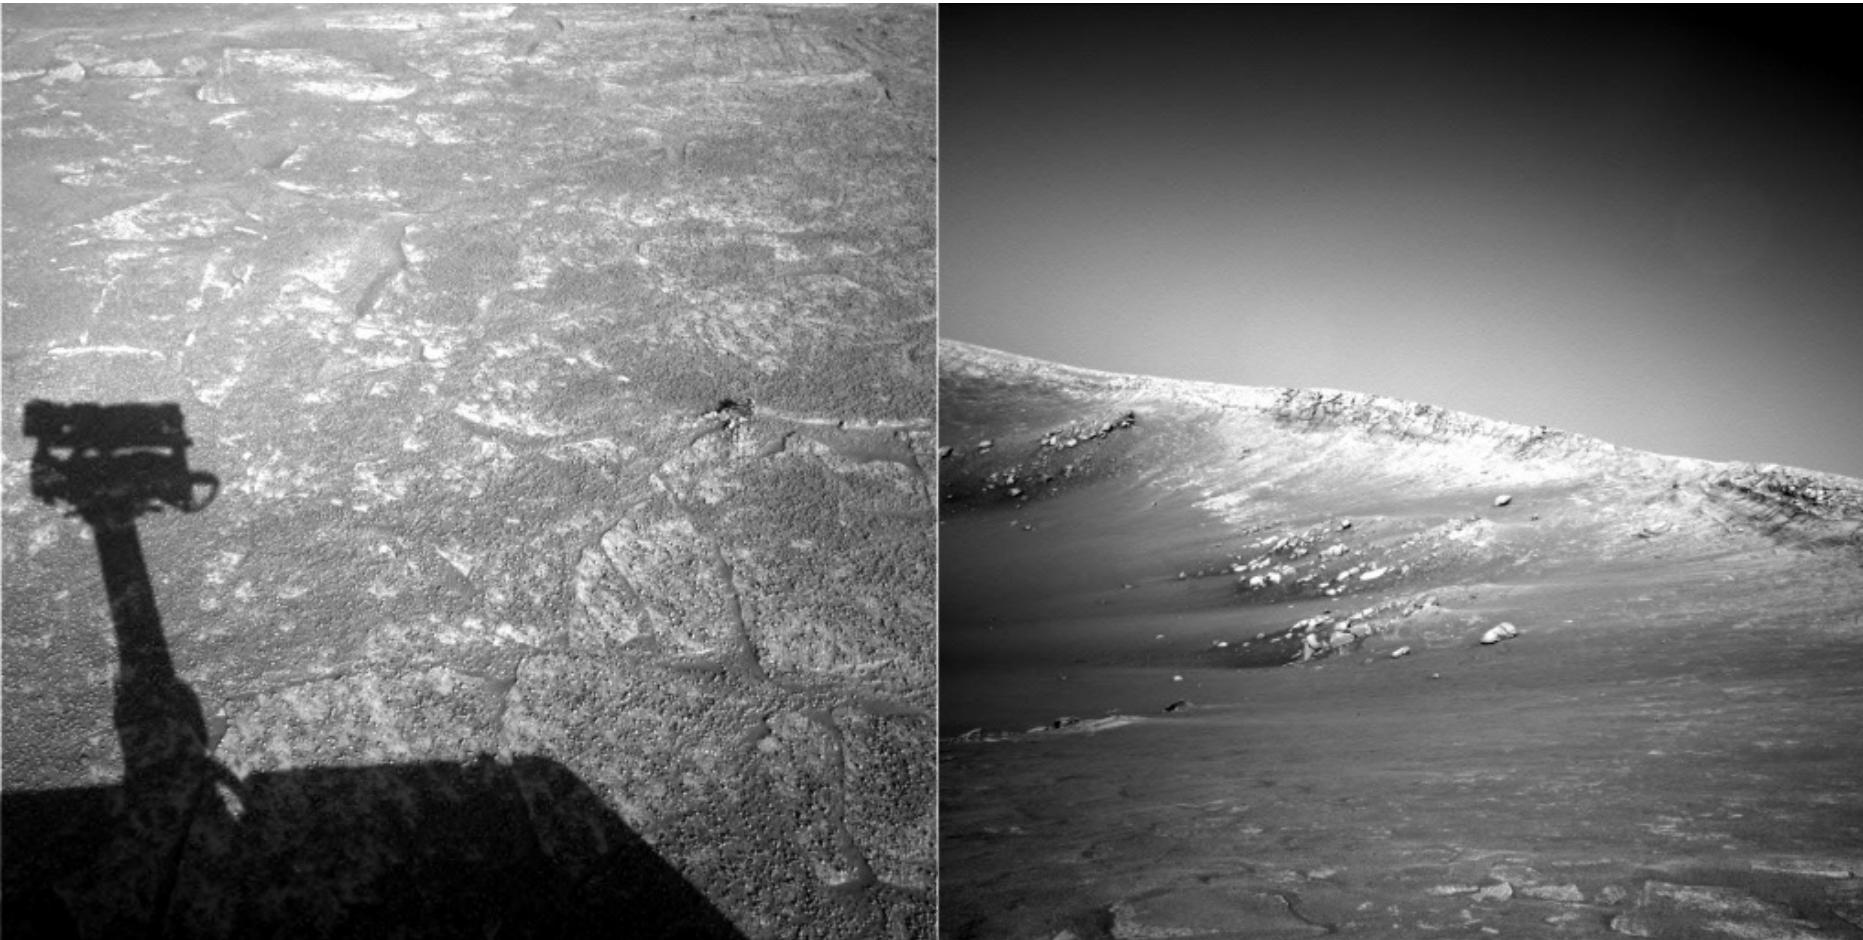


# Robot Localization



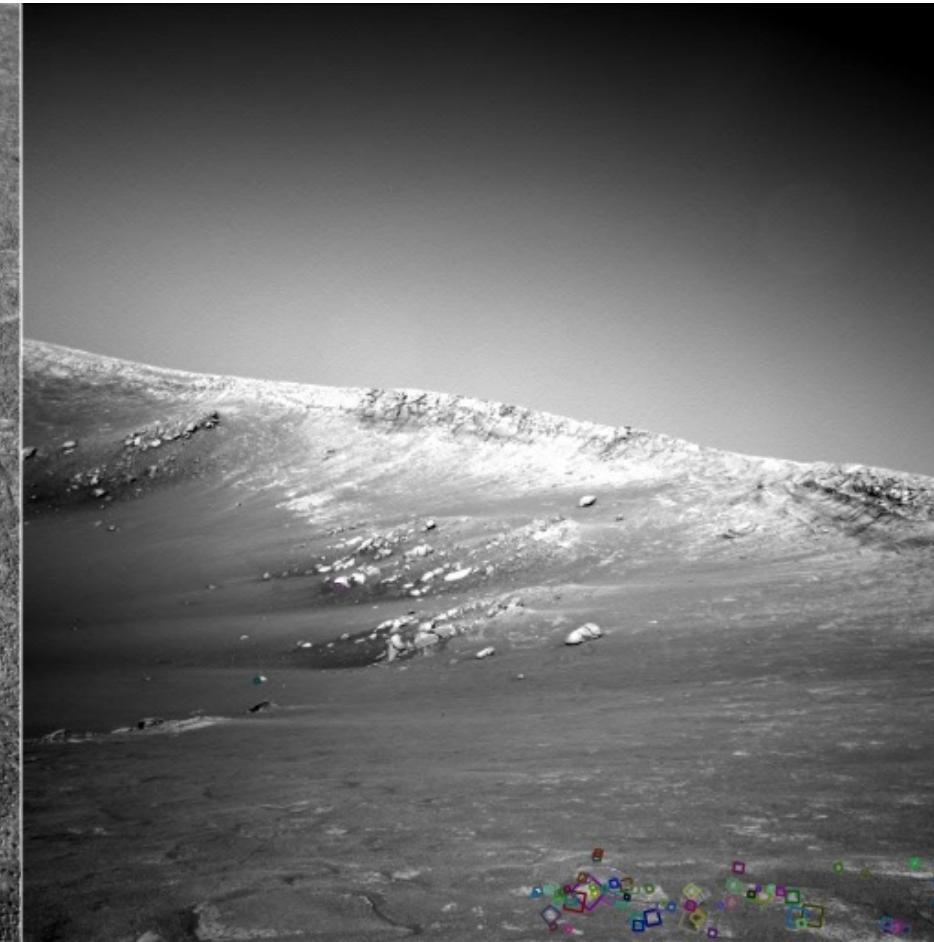
# Image matching





NASA Mars Rover images

*Where are the corresponding points?*



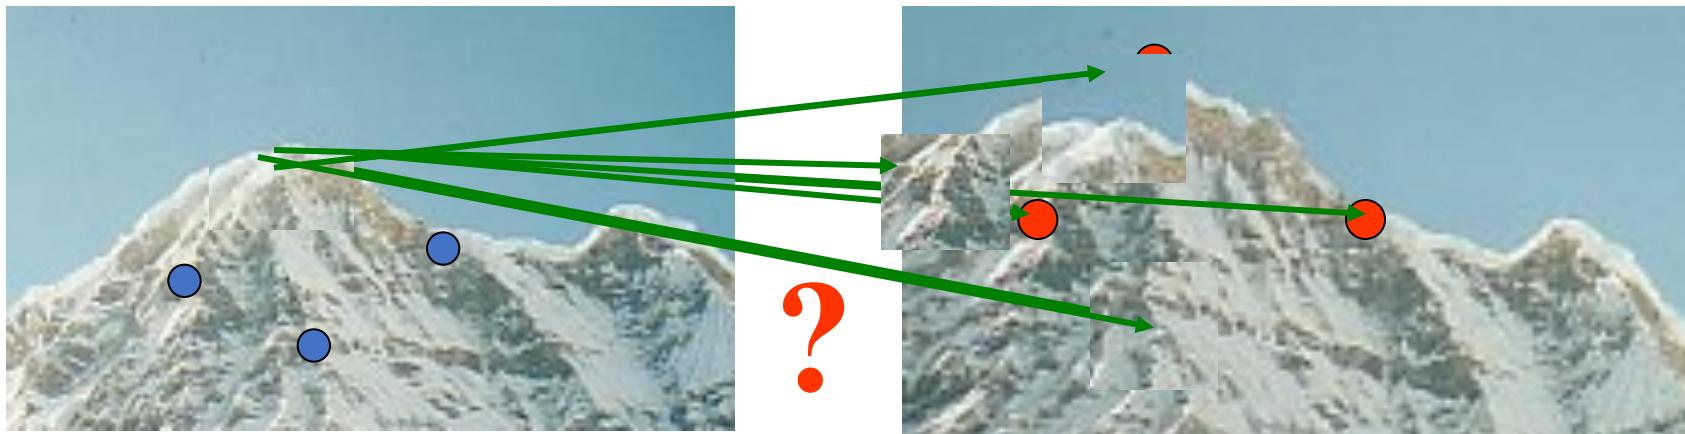
# Example: Panorama Stitching

We have two images –  
how do we overlay them?



# Goal: Distinctiveness

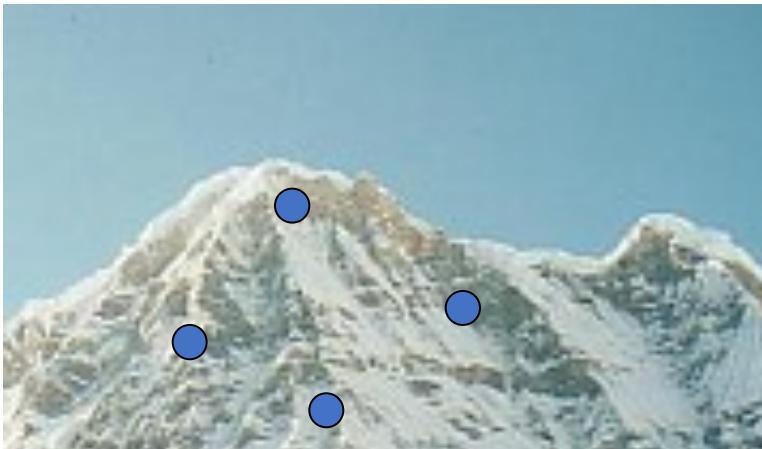
We want to be able to reliably determine which point goes with which.



May be difficult in structured environments  
with repeated elements

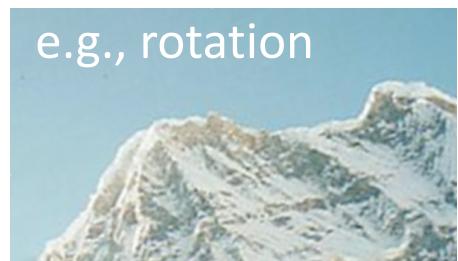
# Goal: Repeatability

We want to detect (at least some of)  
the same points in both images.



With these points, there's no chance to find true matches!

Under geometric and photometric variations.



# Goal: Compactness and Efficiency

We want the representation to be as small and as fast as possible

- Much smaller than a whole image

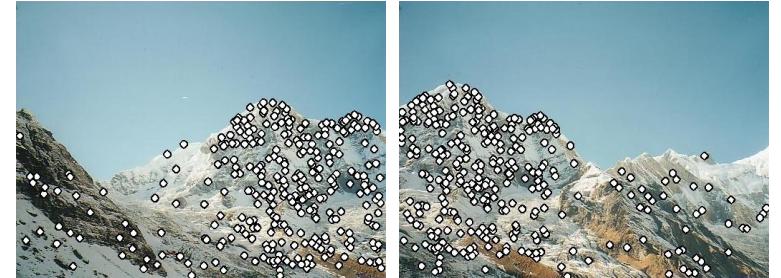
Sometimes, we'd like to run the detection procedure *independently* per image

- Match just the compact descriptors for speed.
- *Difficult!* We don't get to see 'the other image' at match time, e.g., object detection.

# Local Feature Matching: Main Components

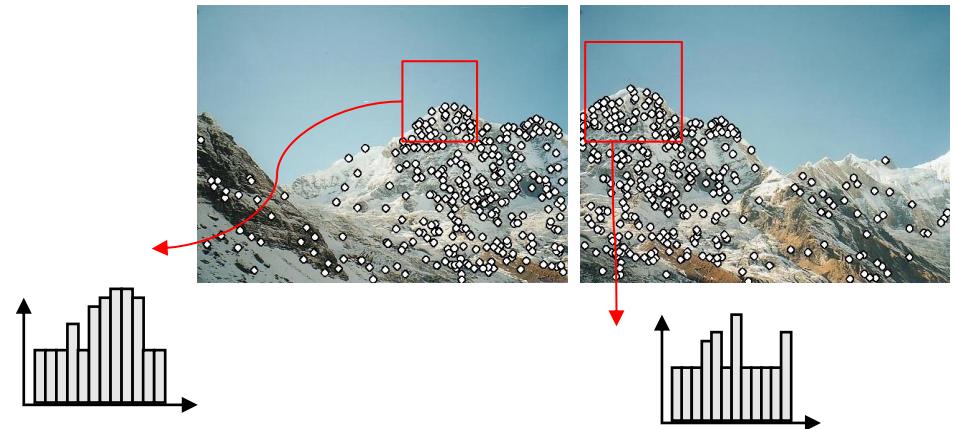
## 1) Detection:

Find a set of distinctive features (e.g., key points)



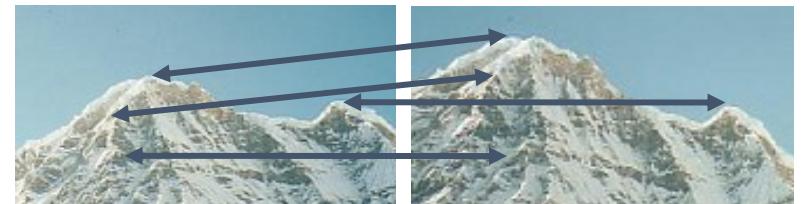
## 2) Description:

Extract feature descriptor around each interest point as vector.



## 3) Matching:

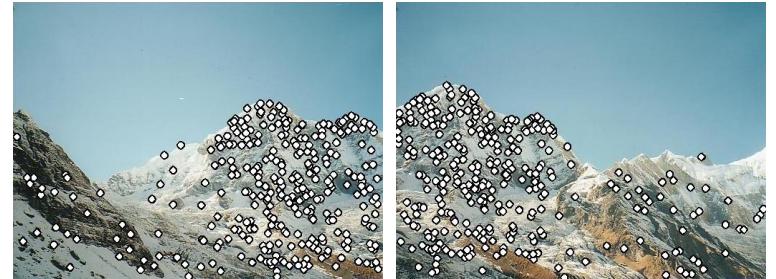
Compute distance between feature vectors to find correspondence.



# Local Feature Matching: Main Components

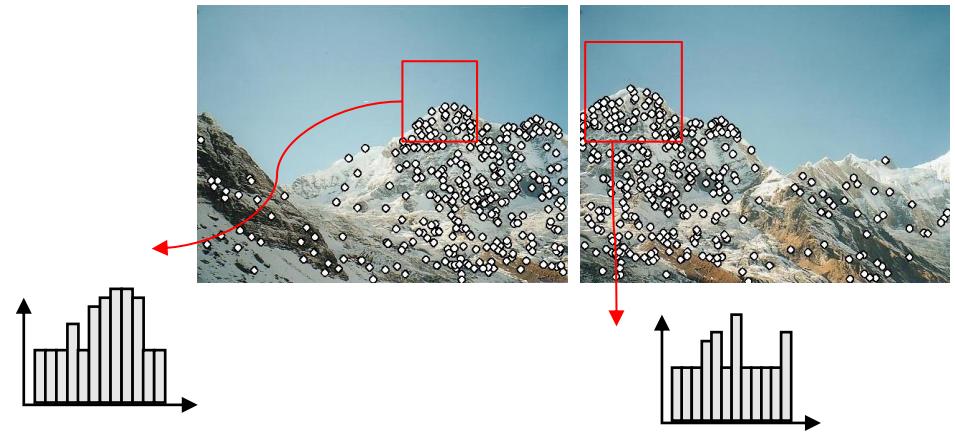
## 1) Detection:

Find a set of distinctive features (e.g., key points)



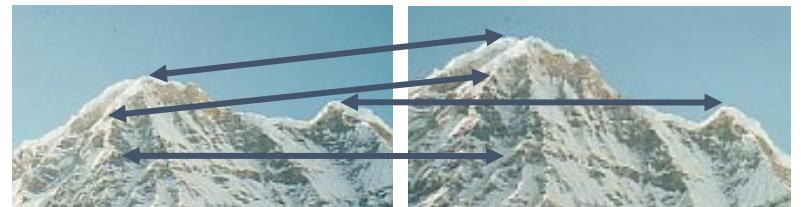
## 2) Description:

Extract feature descriptor around each interest point as vector.



## 3) Matching:

Compute distance between feature vectors to find correspondence.



# Corner Detection: Basic Idea

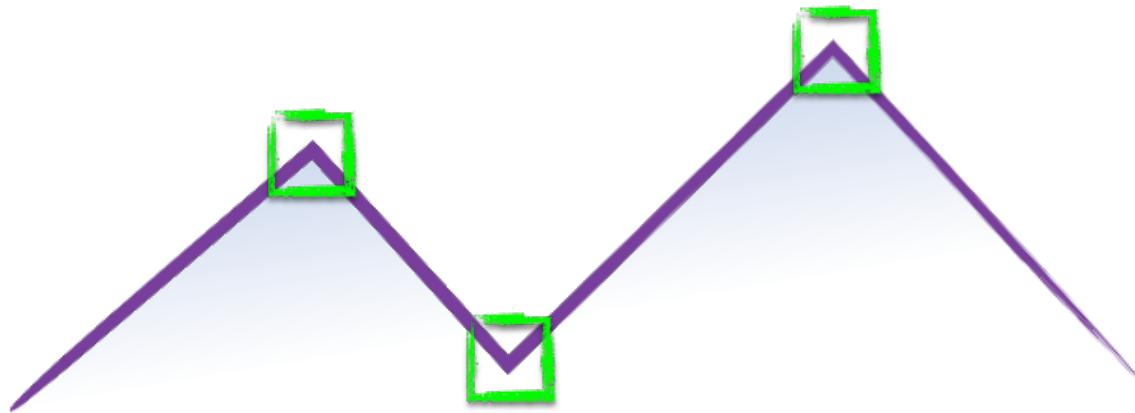
We do not know which other image locations the feature will end up being matched against.

*But we can compute how stable a location is in appearance with respect to small variations in its position.*

*Strategy: Compare image patch against local neighbors.*

# How do you find a corner?

[Moravec 1980]

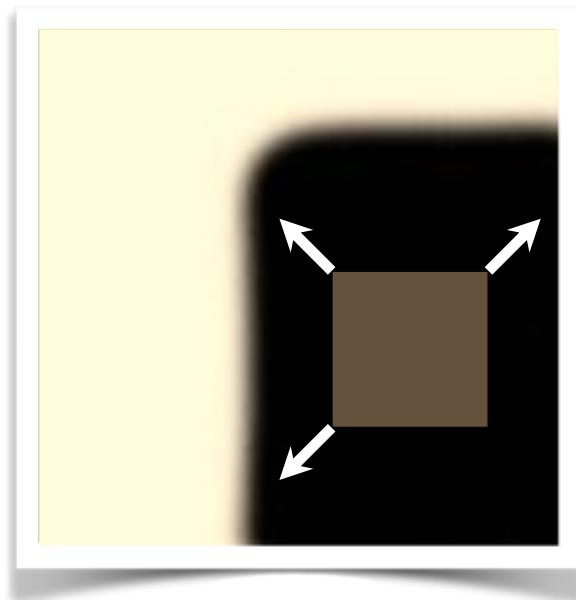


Easily recognized by looking through a small window

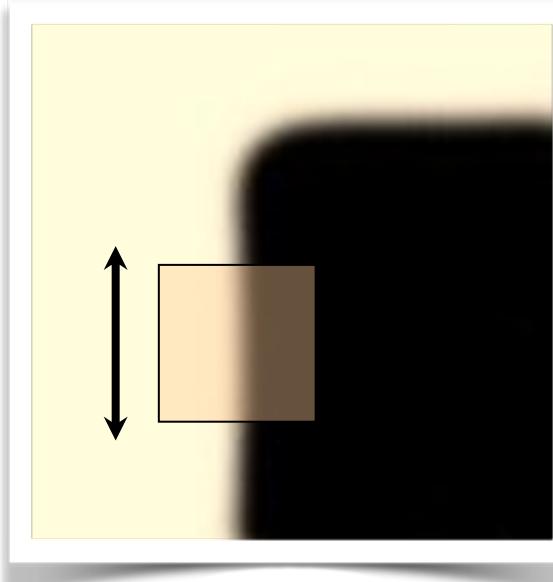
Shifting the window should give large change in intensity

Easily recognized by looking through a small window

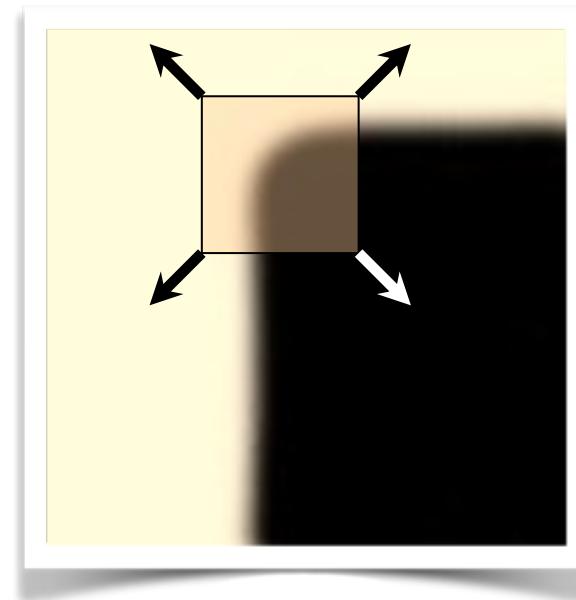
Shifting the window should give large change in intensity



“flat” region:  
no change in all  
directions



“edge”:  
no change along the edge  
direction



“corner”:  
significant change in all  
directions

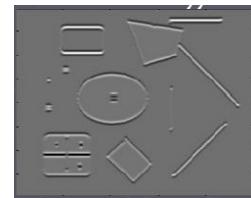
# Finding corners

1. Compute image gradients over small region

$$I_x = \frac{\partial I}{\partial x}$$

$$I_y = \frac{\partial I}{\partial y}$$

2. Subtract mean from each image gradient



3. Compute the covariance matrix

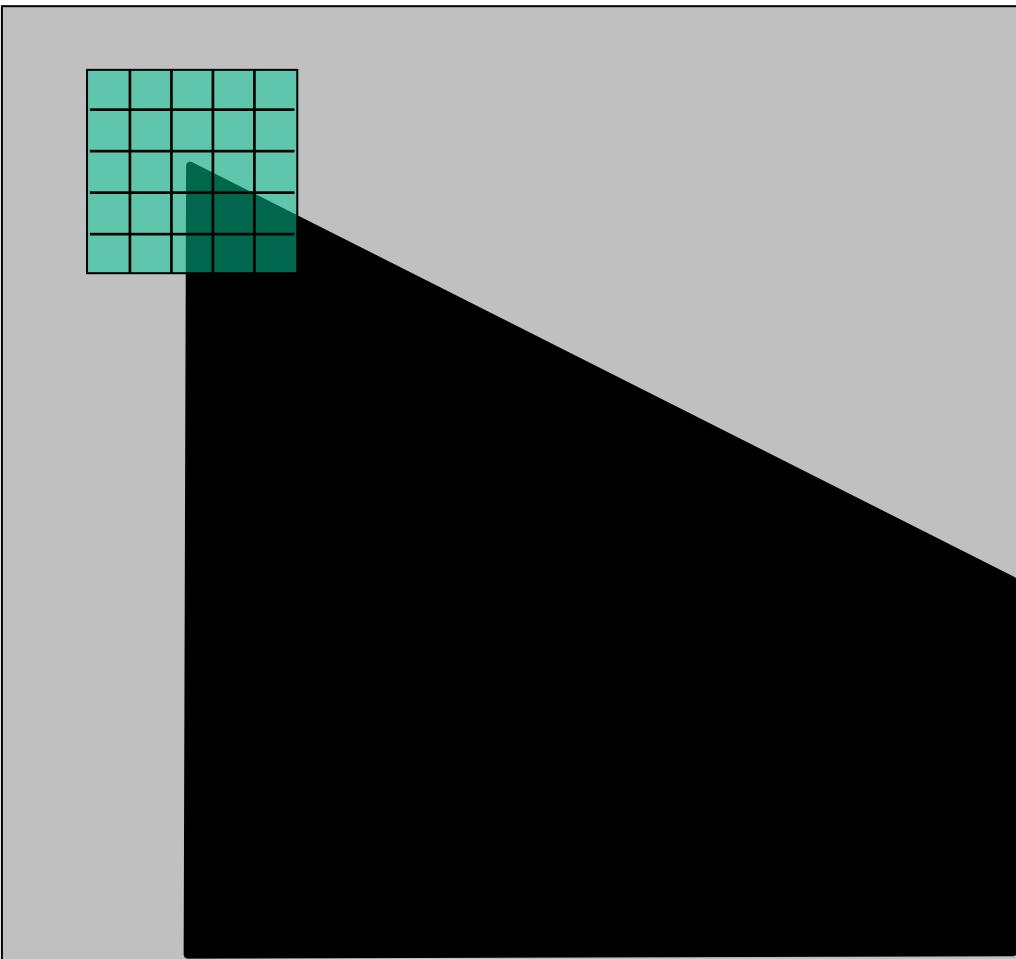
4. Compute eigenvectors and eigenvalues

$$\begin{bmatrix} \sum_{p \in P} I_x I_x & \sum_{p \in P} I_x I_y \\ \sum_{p \in P} I_y I_x & \sum_{p \in P} I_y I_y \end{bmatrix}$$

5. Use threshold on eigenvalues to detect corners

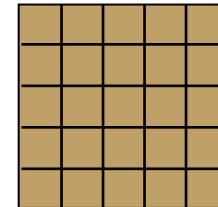
1. Compute image gradients over a small region  
(not just a single pixel)

# 1. Compute image gradients over a small region (not just a single pixel)



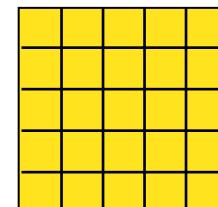
array of x gradients

$$I_x = \frac{\partial I}{\partial x}$$



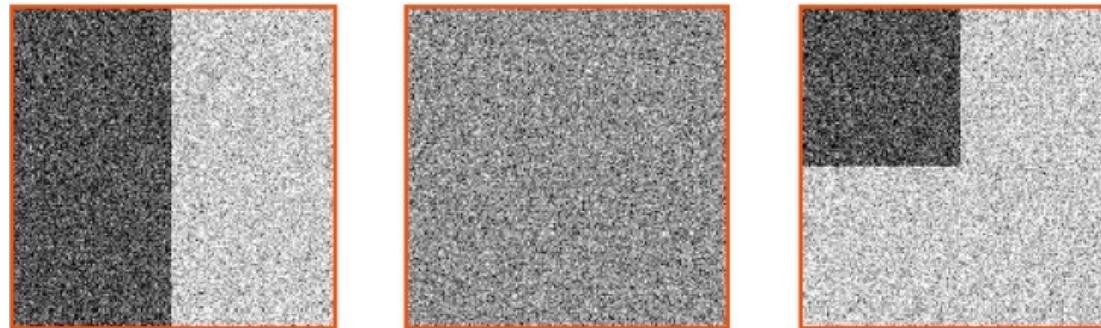
array of y gradients

$$I_y = \frac{\partial I}{\partial y}$$

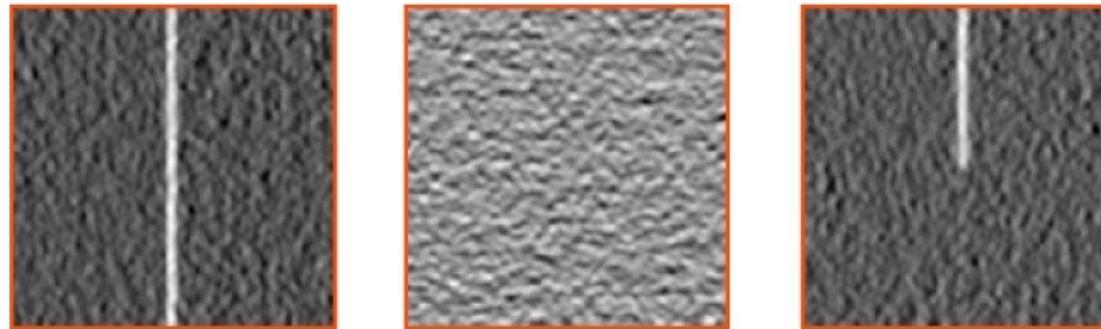


# visualization of gradients

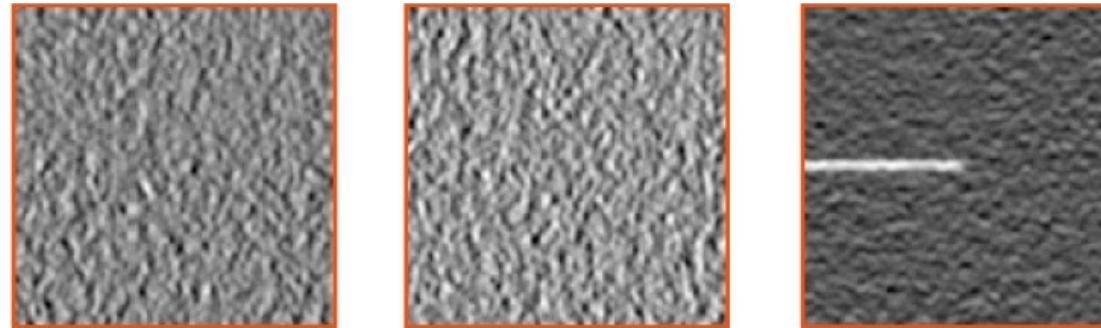
image

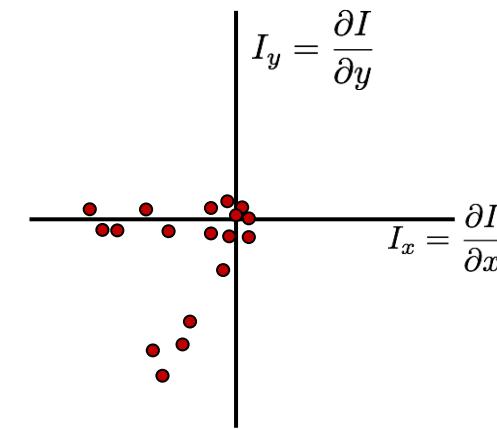
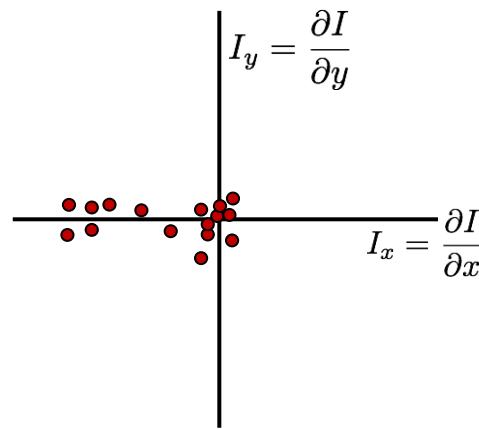
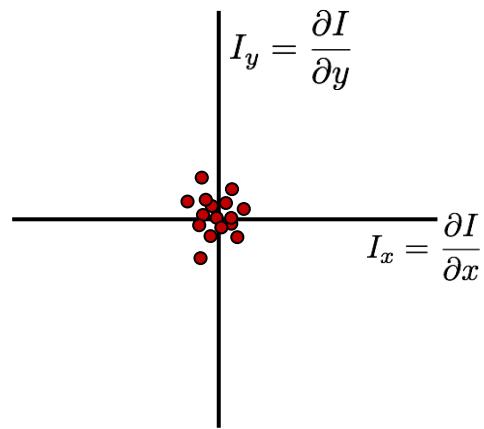
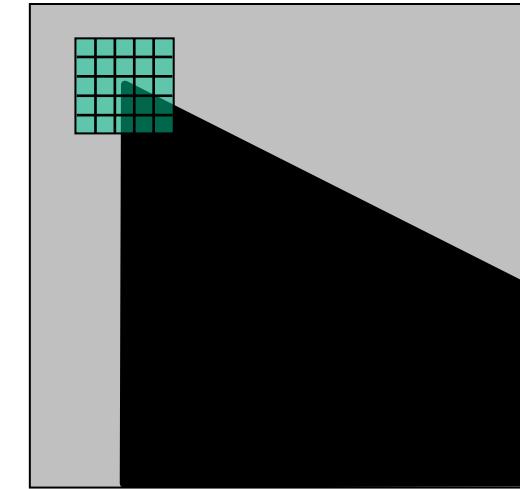
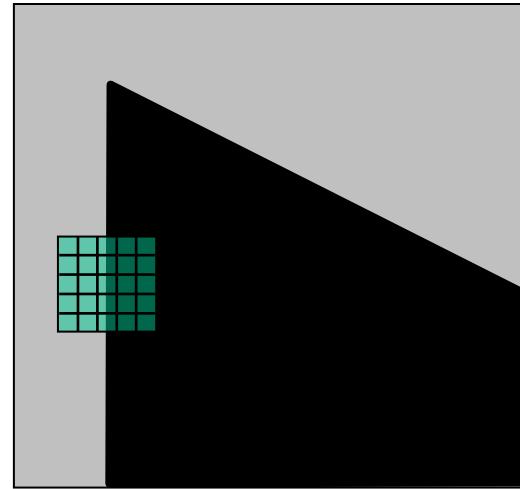
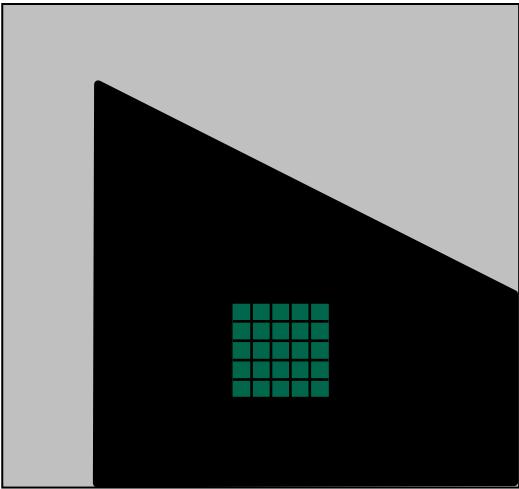


X derivative

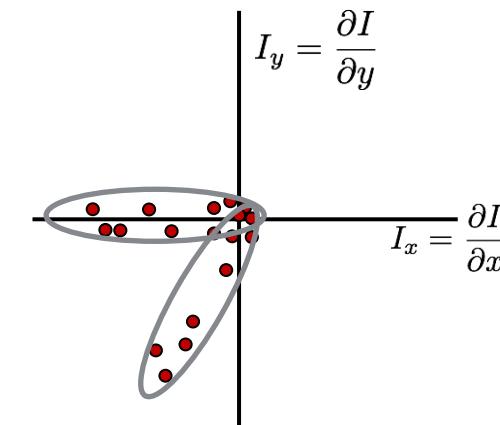
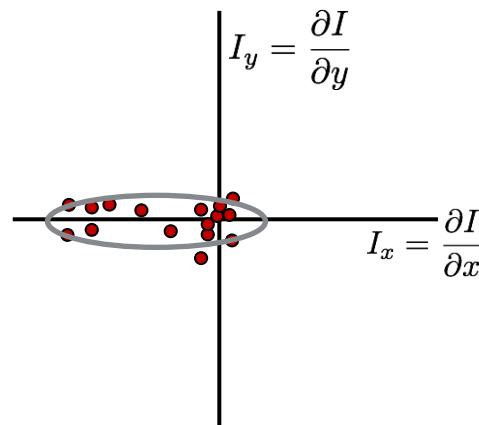
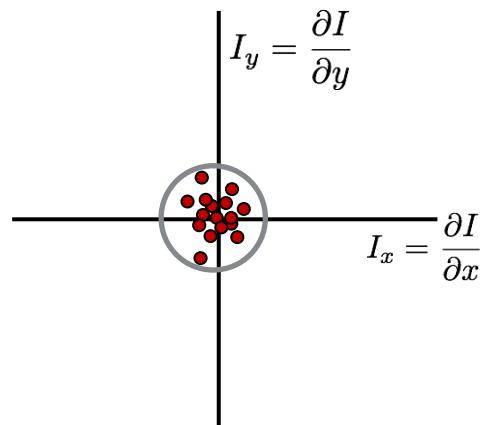
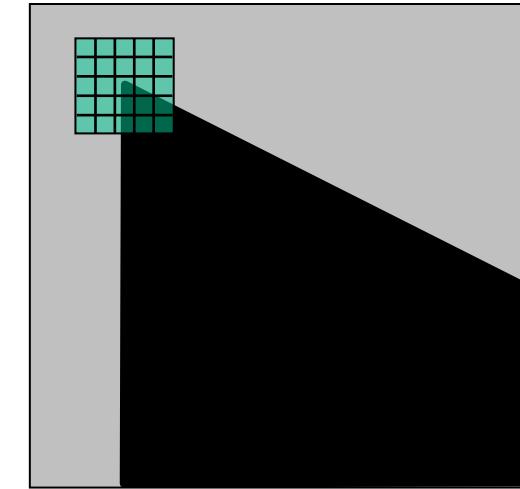
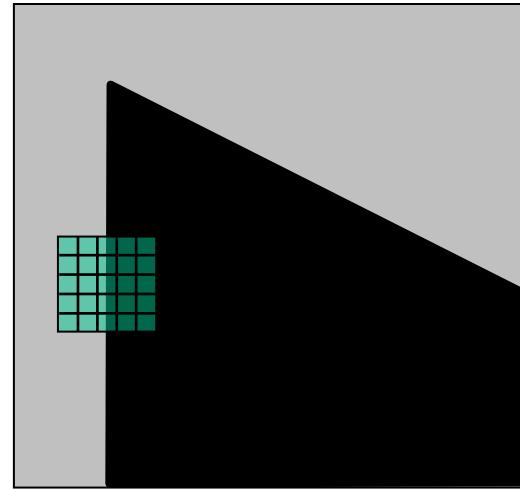
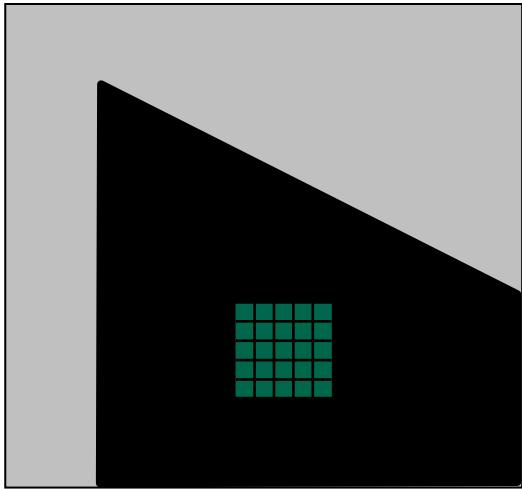


Y derivative

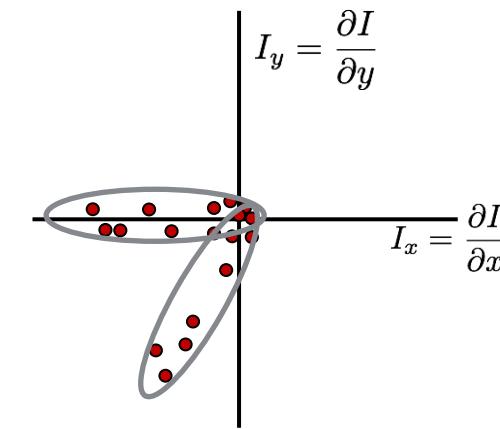
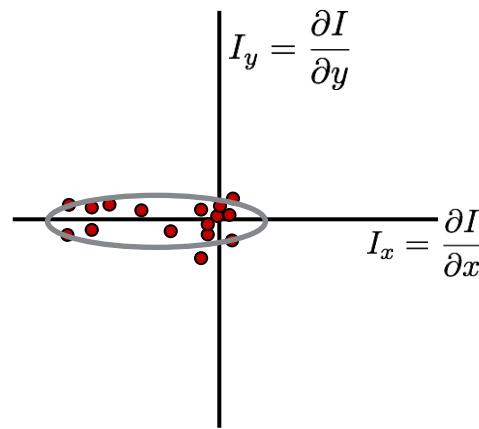
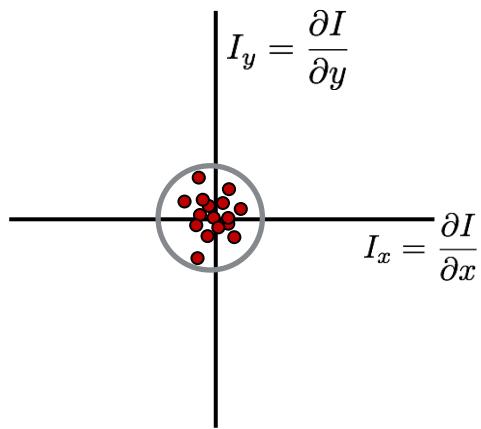
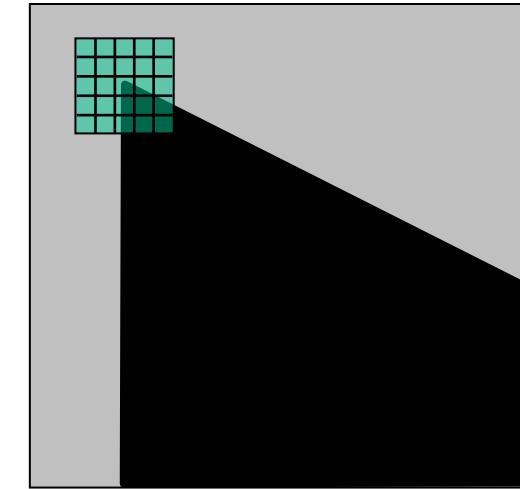
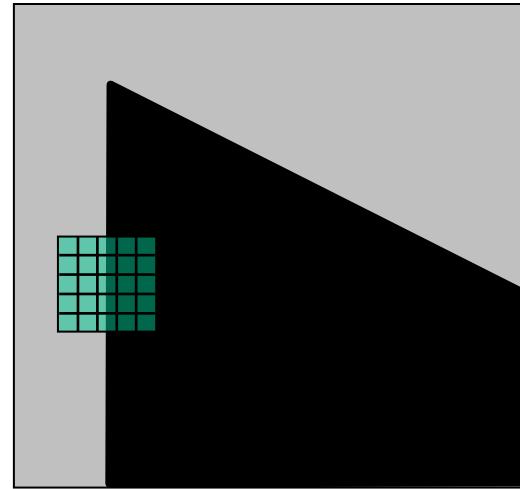
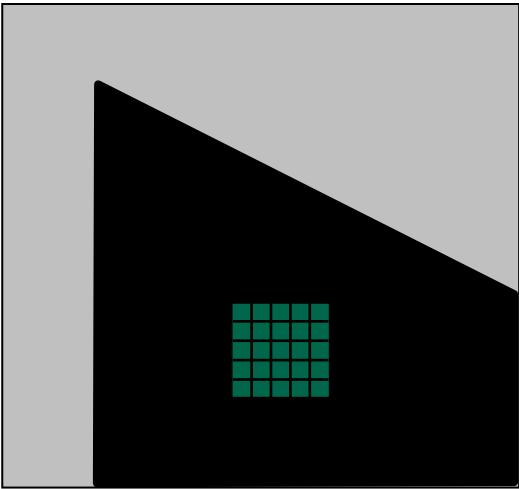




*What does the distribution tell you about the region?*



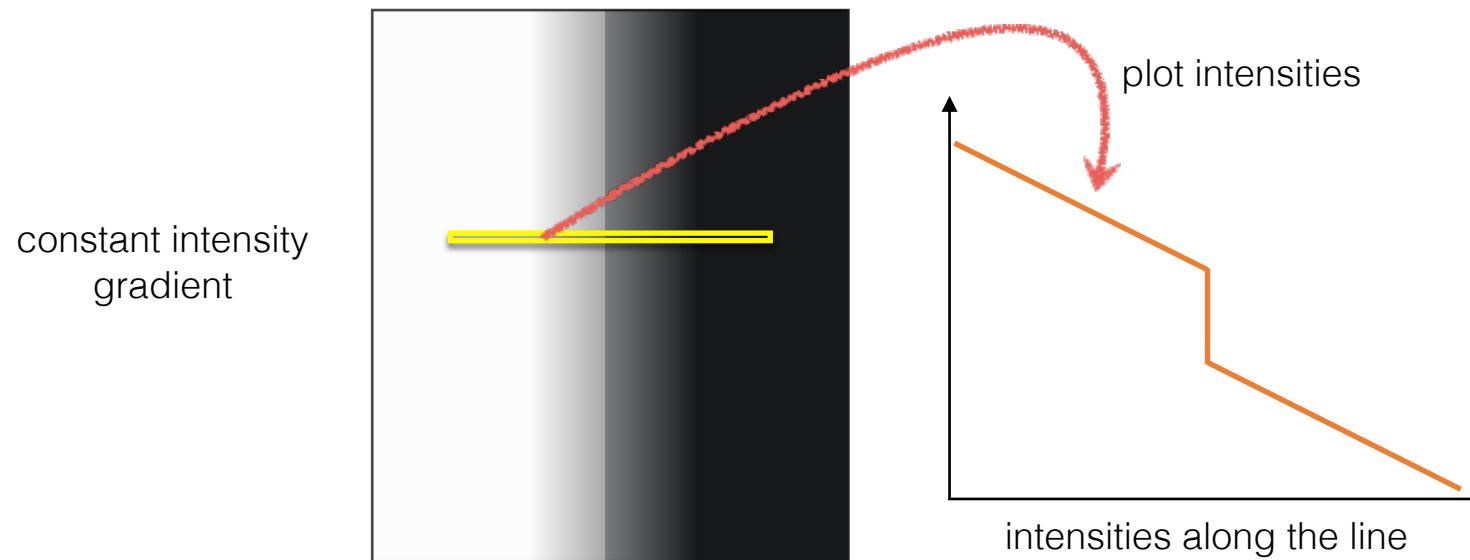
distribution reveals edge orientation and magnitude



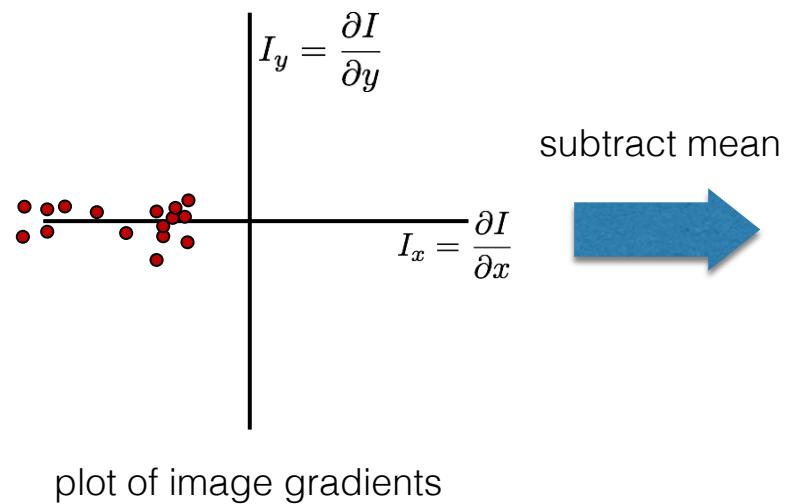
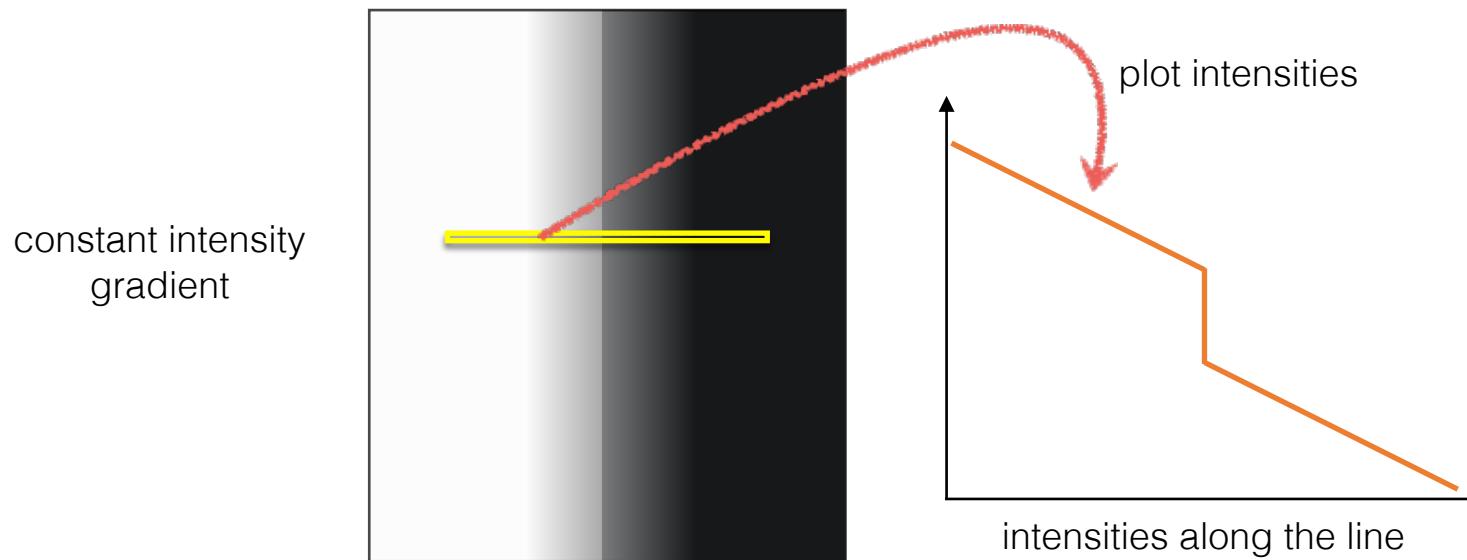
*How do you quantify orientation and magnitude?*

2. Subtract the mean from each image gradient

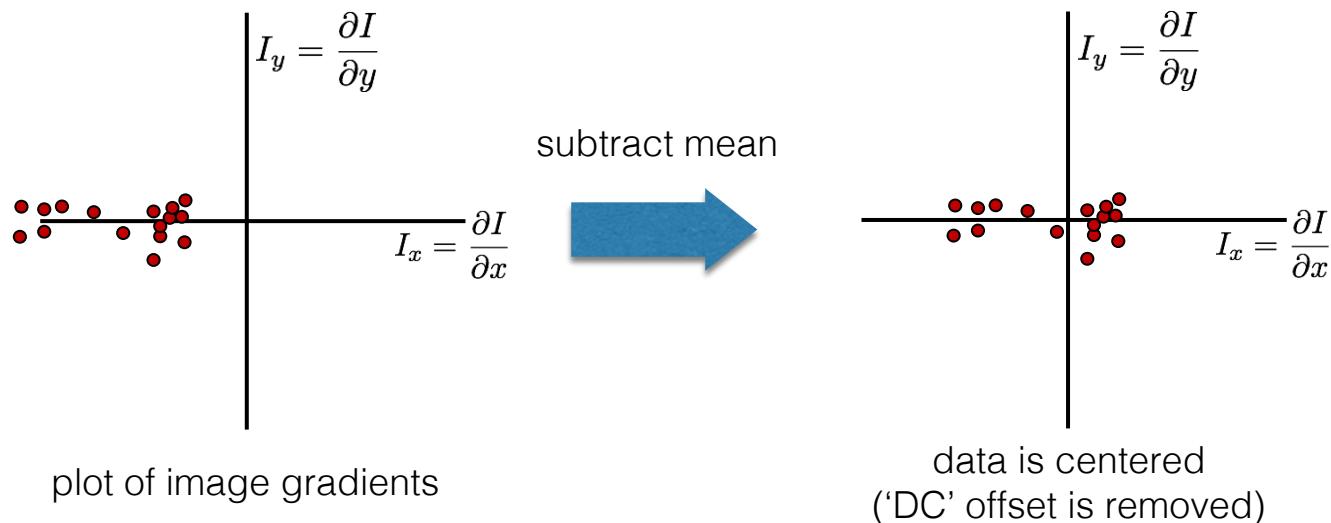
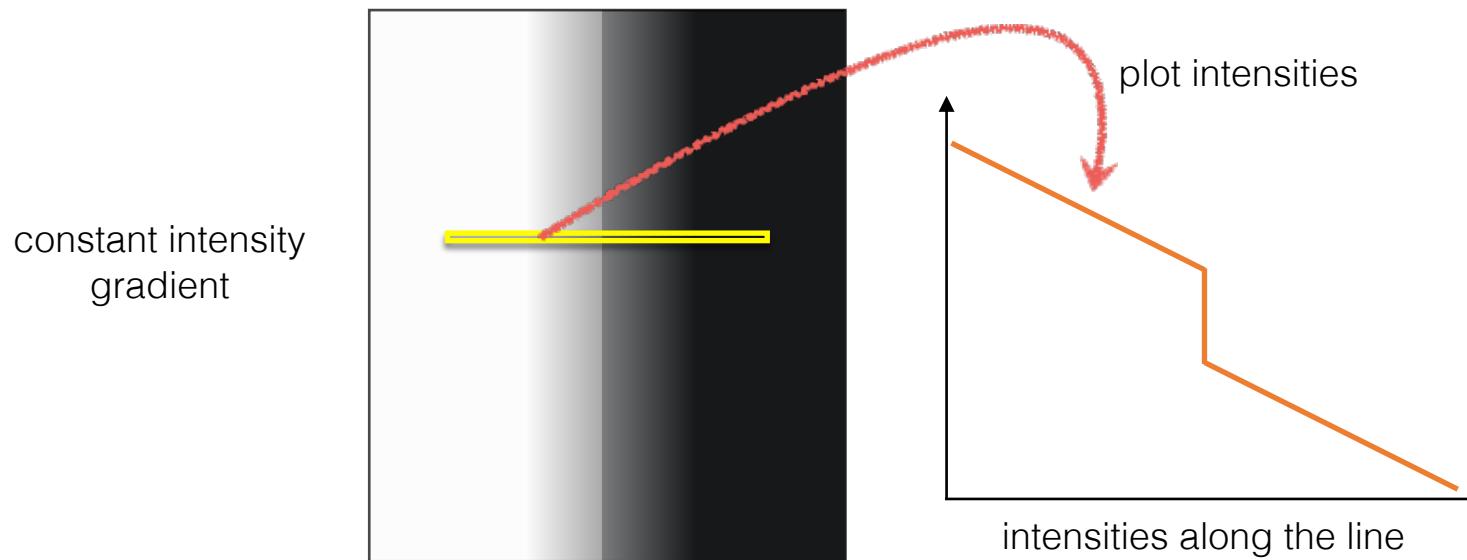
## 2. Subtract the mean from each image gradient



## 2. Subtract the mean from each image gradient



## 2. Subtract the mean from each image gradient



3. Compute the covariance matrix

### 3. Compute the covariance matrix

$$\begin{bmatrix} \sum_{p \in P} I_x I_x & \sum_{p \in P} I_x I_y \\ \sum_{p \in P} I_y I_x & \sum_{p \in P} I_y I_y \end{bmatrix}$$

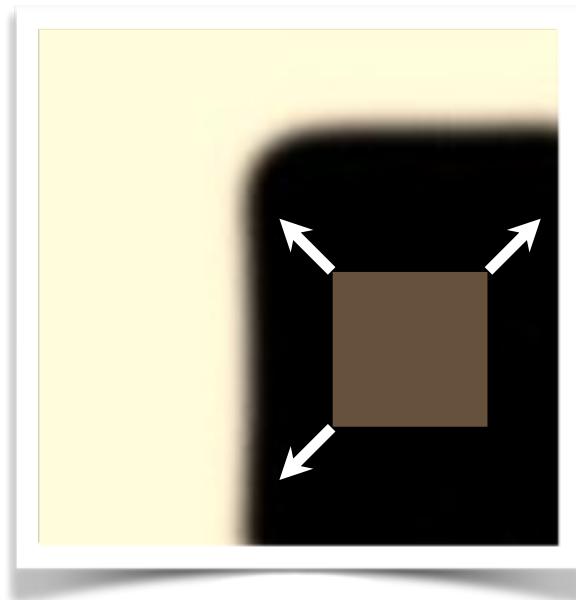
$$\sum_{p \in P} I_x I_y = \text{sum}\left( \begin{array}{|c|c|c|c|} \hline & & & \\ \hline \end{array} \right) \cdot \begin{array}{|c|c|c|c|} \hline & & & \\ \hline \end{array} )$$

array of x gradients                          array of y gradients

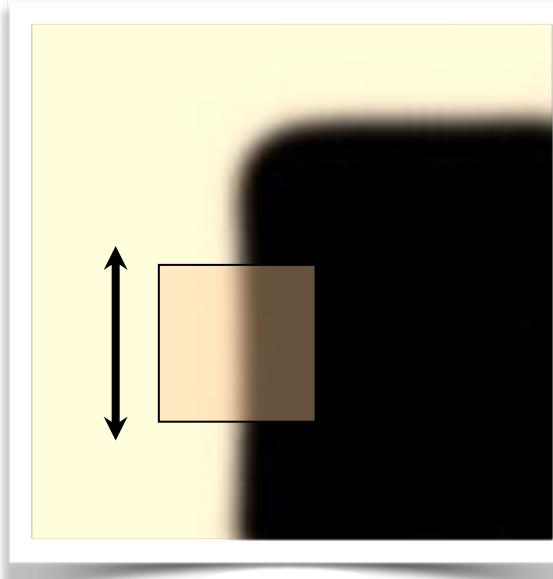
*Where does this covariance matrix come from?*

Easily recognized by looking through a small window

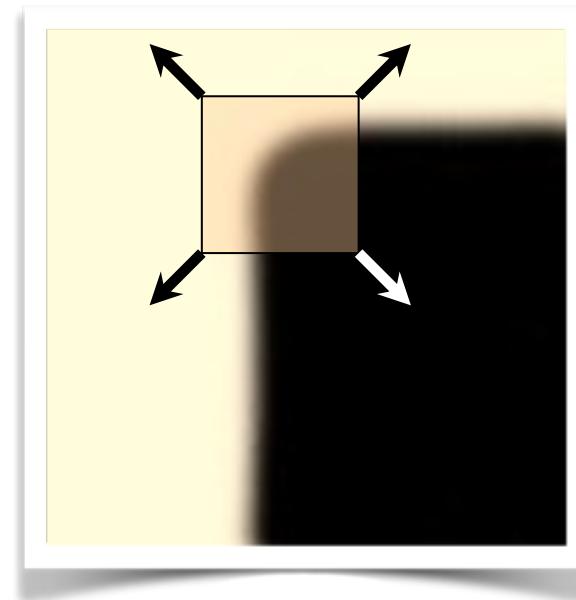
Shifting the window should give large change in intensity



“flat” region:  
no change in all  
directions



“edge”:  
no change along the edge  
direction

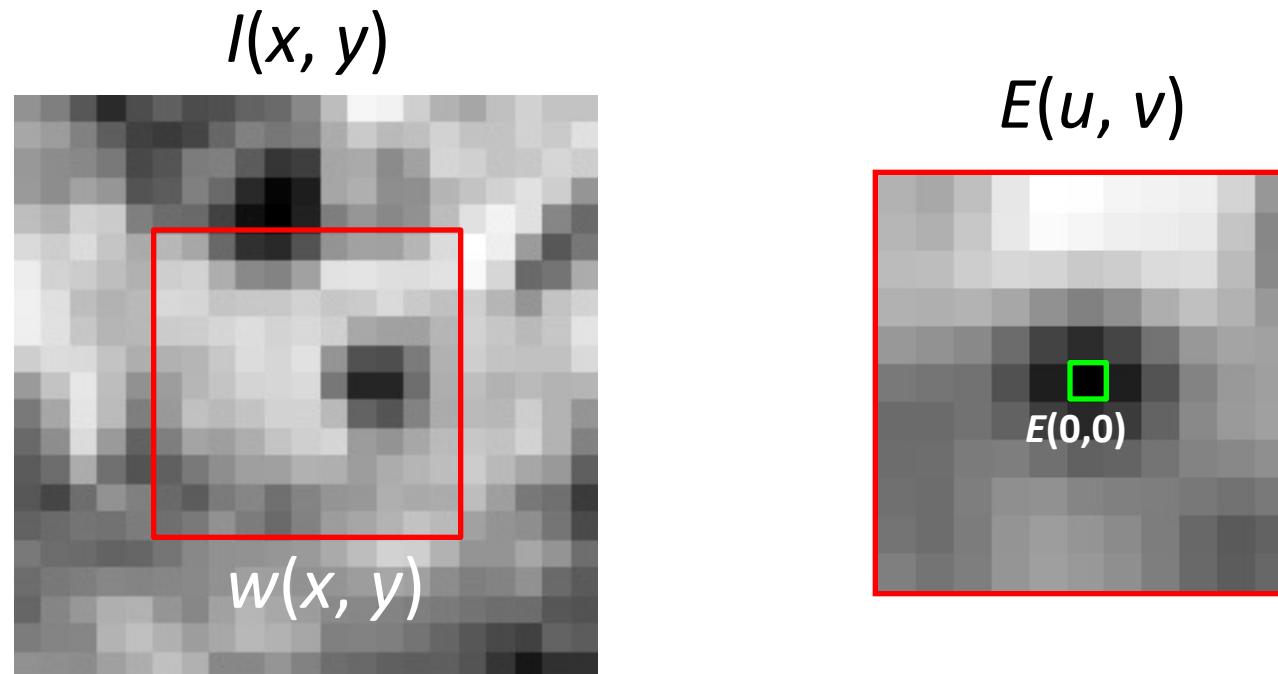


“corner”:  
significant change in all  
directions

# Corner Detection by Auto-correlation

Change in appearance of window  $w(x,y)$  for shift  $[u,v]$ :

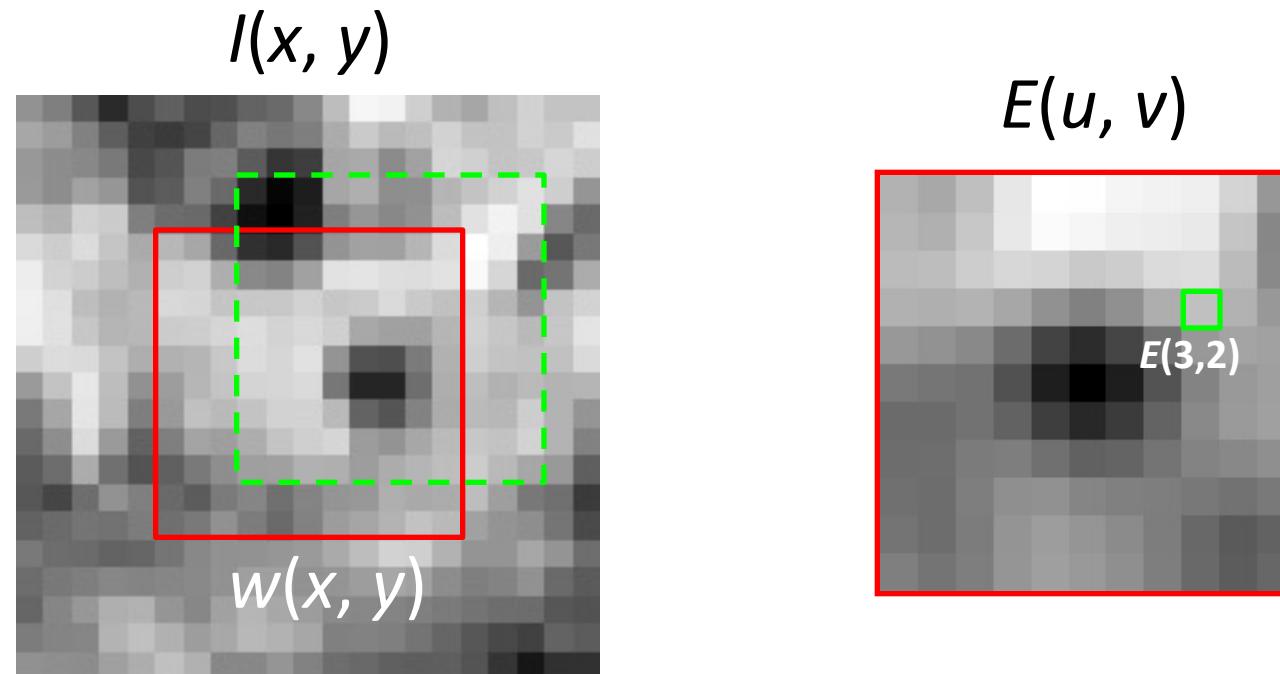
$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$



# Corner Detection by Auto-correlation

Change in appearance of window  $w(x,y)$  for shift  $[u,v]$ :

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$



# Visualizing Error function

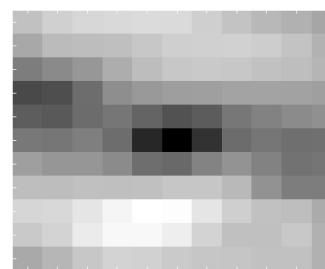
$$E(u, v) = \sum_{x, y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

Correspond the three red crosses to (b,c,d).

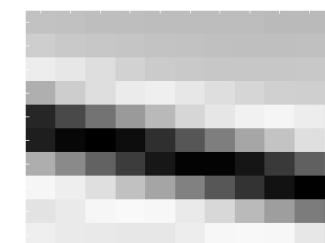


$E(u, v)$

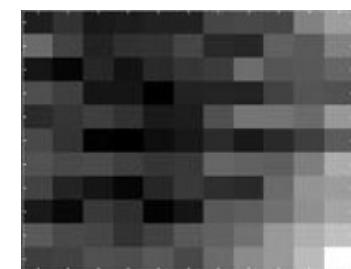
b)



c)

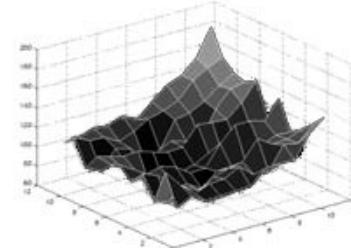
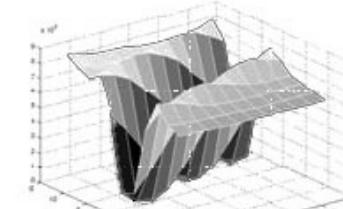
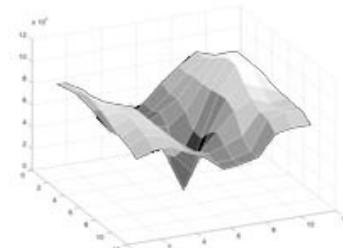


d)



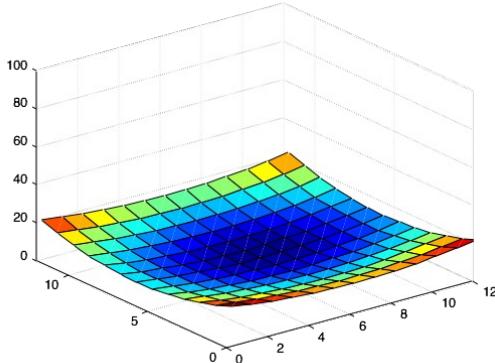
$E(u, v)$

As a surface

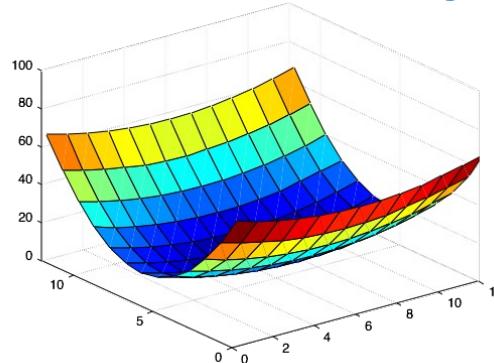


- b) is from the flower bed (good unique minimum);
- c) is from the roof edge (one-dimensional aperture problem); and
- d) is from the cloud (no good peak).

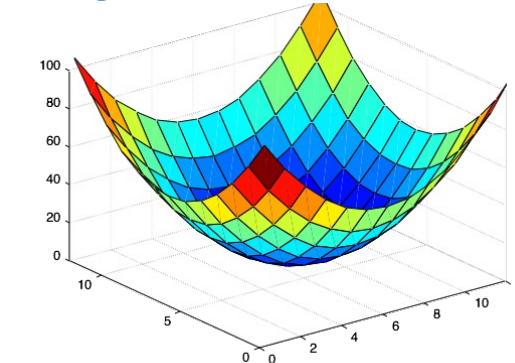
## *Which error surface indicates a good image feature?*



flat



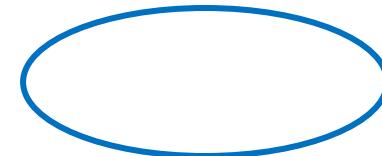
edge  
'line'



corner  
'dot'

*What kind of image patch do these surfaces represent?*

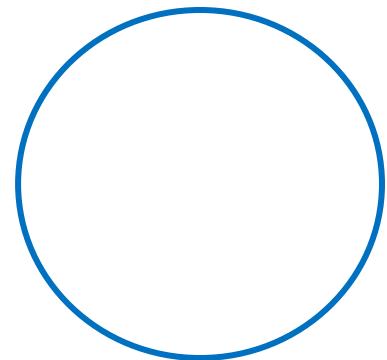
Assume, we plot for  
 $f(x,y) = \text{some fixed}$   
value  $p$ .



edge  
'line'



corner  
'dot'



flat

# Corner Detection by Auto-correlation

Without loss of generality, assume a grayscale 2-dimensional image,  $I$  is used.  
Consider taking an image patch  $(x, y) \in W$  (window) and shift it by  $(\Delta x, \Delta y)$ .

The sum of squared differences (SSD) between these two patches,  $f$  is :

$$f(\Delta x, \Delta y) = \sum_{(x_k, y_k) \in W} (I(x_k, y_k) - I(x_k + \Delta x, y_k + \Delta y))^2$$

$I(x + \Delta x, y + \Delta y)$  can be approximated by a Taylor expansion.

Let  $I_x$  and  $I_y$  be the partial derivatives of  $I$ , such that

$$I(x + \Delta x, y + \Delta y) \approx I(x, y) + I_x(x, y)\Delta x + I_y(x, y)\Delta y$$

$$\Rightarrow f(\Delta x, \Delta y) \approx \sum_{(x, y) \in W} (I_x(x, y)\Delta x + I_y(x, y)\Delta y)^2$$

Alternatively,  $f(\Delta x, \Delta y) \approx (\Delta x \quad \Delta y) M \begin{pmatrix} \Delta x \\ \Delta y \end{pmatrix}$

where  $M$  is the structure tensor, given by :

$$M = \sum_{(x, y) \in W} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} = \begin{bmatrix} \sum_{(x, y) \in W} I_x^2 & \sum_{(x, y) \in W} I_x I_y \\ \sum_{(x, y) \in W} I_x I_y & \sum_{(x, y) \in W} I_y^2 \end{bmatrix}$$

# Error function

Change of intensity for the shift  $[u, v]$ :

$$E(u, v) = \sum_{x,y} w(x, y) [I(x+u, y+v) - I(x, y)]^2 \rightarrow E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

↑      ↑      ↑      ↑  
Error   Window   Shifted   Intensity  
function   function   intensity

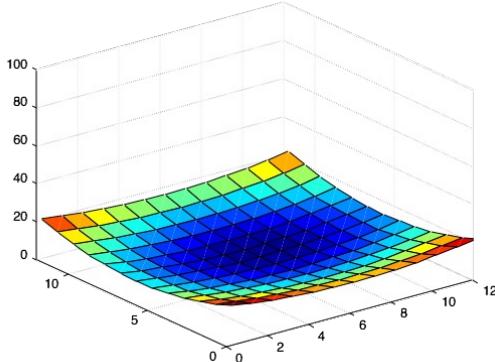
Window function  $w(x, y) =$

1 in window, 0 outside      or      Gaussian

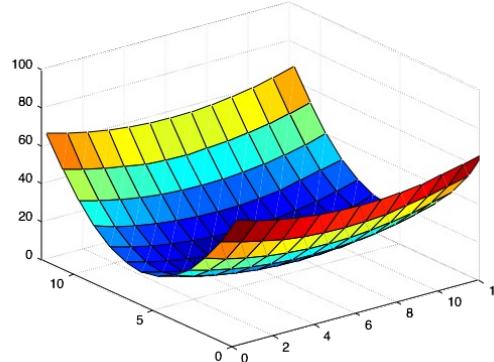
$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

'second moment' matrix  
'structure tensor'

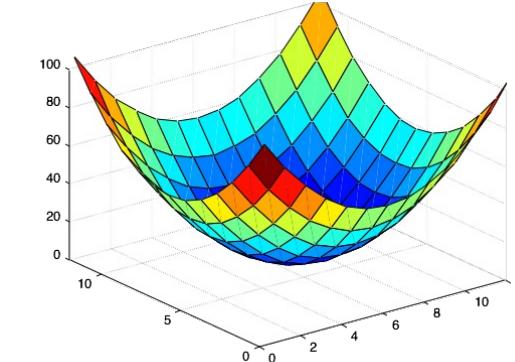
## *Which error surface indicates a good image feature?*



flat



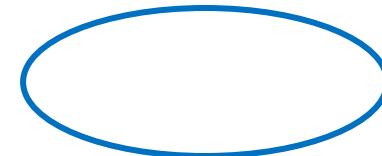
edge  
'line'



corner  
'dot'

*What kind of image patch do these surfaces represent?*

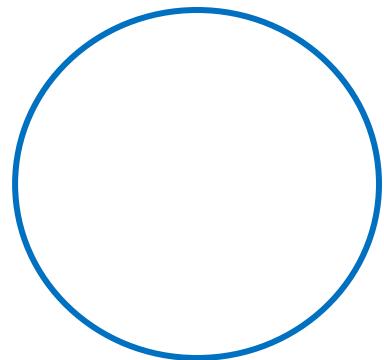
Assume, we plot for  
 $f(x,y) = \text{some fixed}$   
value  $p$ .



edge  
'line'



corner  
'dot'



flat

4. Compute eigenvalues and eigenvectors

#### 4. Compute eigenvalues and eigenvectors

$$M\mathbf{e} = \lambda\mathbf{e}$$

↑  
eigenvalue  
↑    ↑  
 $M - \lambda I$        $\mathbf{e} = 0$   
eigenvector

1. Compute the determinant of  
(returns a polynomial)

$$M - \lambda I$$

2. Find the roots of polynomial  
(returns eigenvalues)

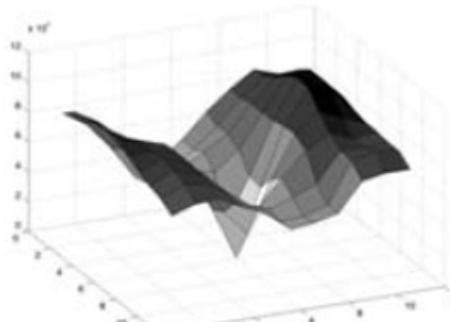
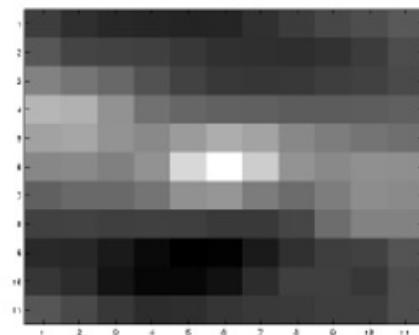
$$\det(M - \lambda I) = 0$$

3. For each eigenvalue, solve  
(returns eigenvectors)

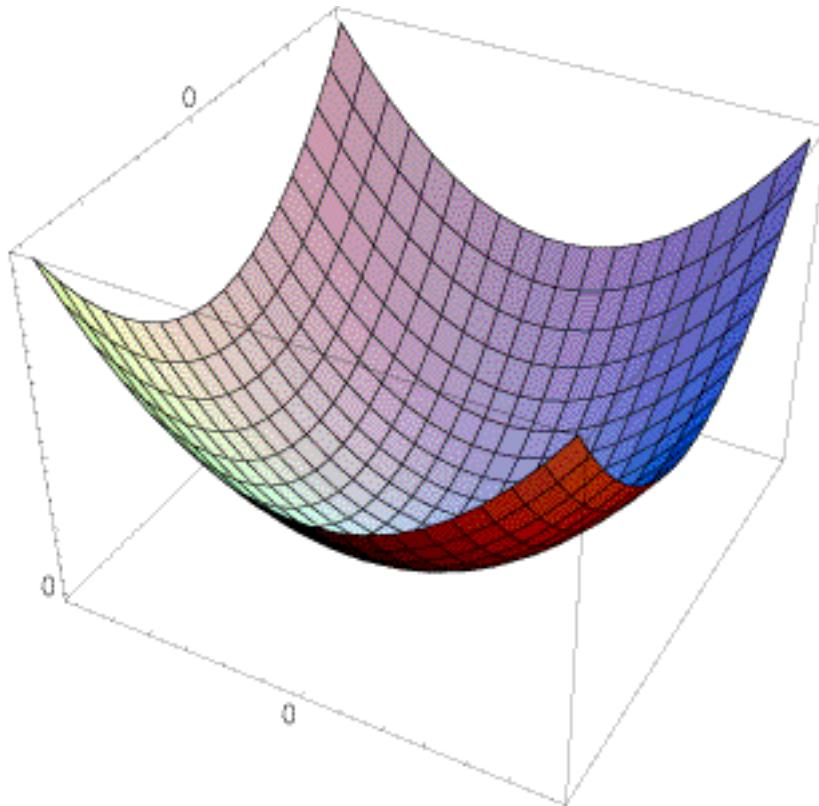
$$(M - \lambda I)\mathbf{e} = 0$$

# Strategy

Approximate  $E(u,v)$  locally by a quadratic surface, and look for that instead.



$\approx$



For cornerness, we only care about the ‘steepness’, not the rotation.

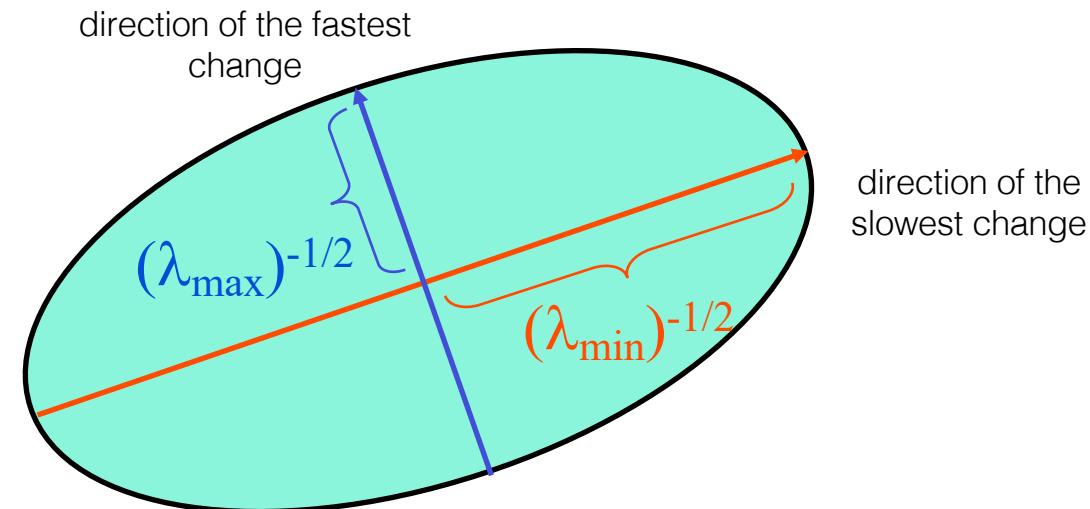
# Visualization as an ellipse

Since  $M$  is symmetric, we have  $M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$

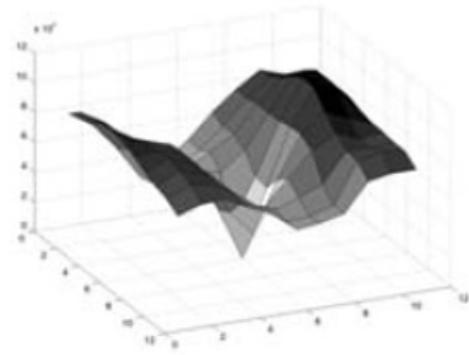
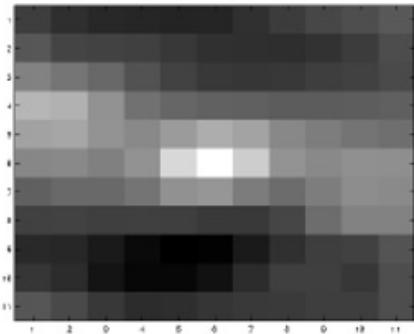
We can visualize  $M$  as an ellipse with axis lengths determined by the eigenvalues and orientation determined by  $R$

Ellipse equation:

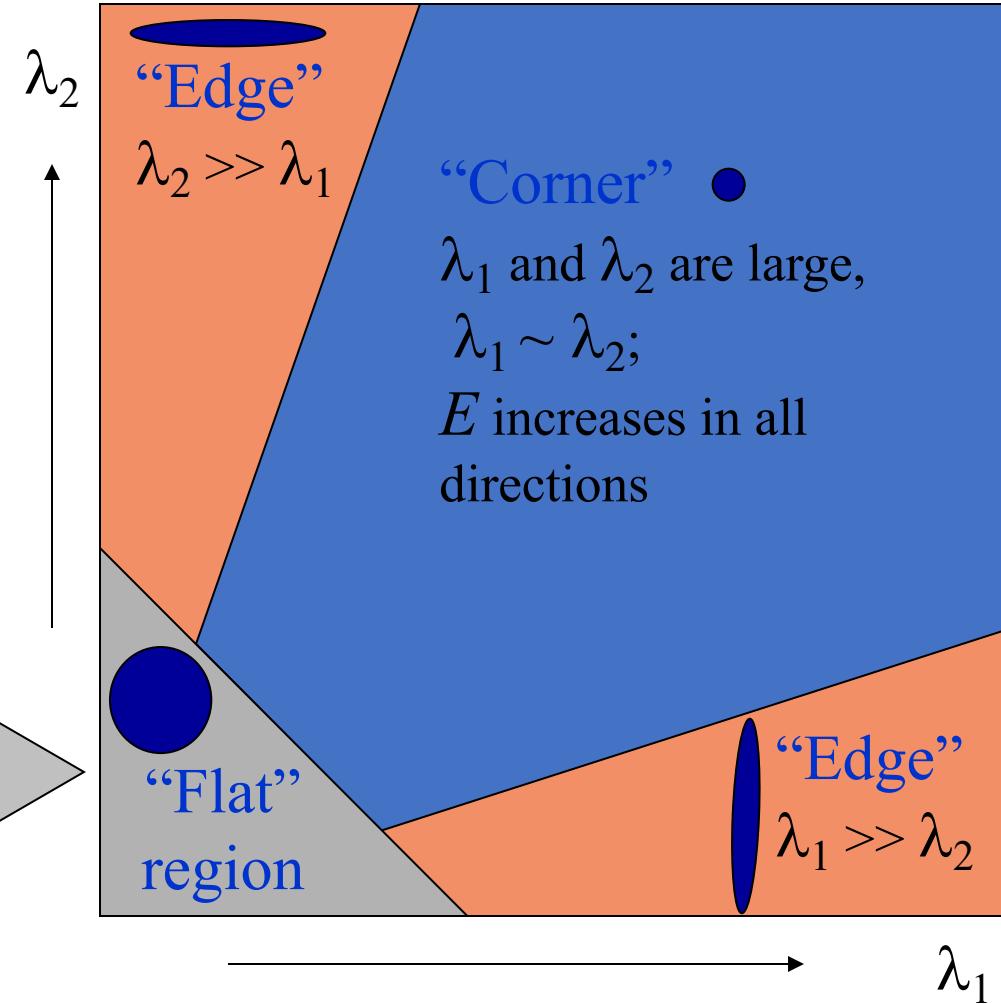
$$[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$



# Classification of image points using eigenvalues of $M$



$\lambda_1$  and  $\lambda_2$  are small;  
 $E$  is almost constant  
in all directions

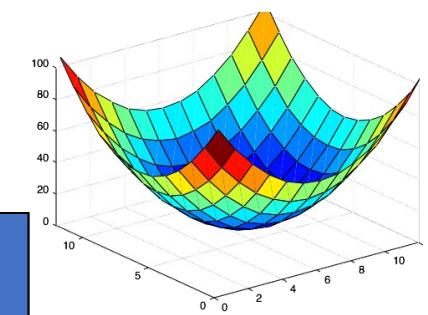
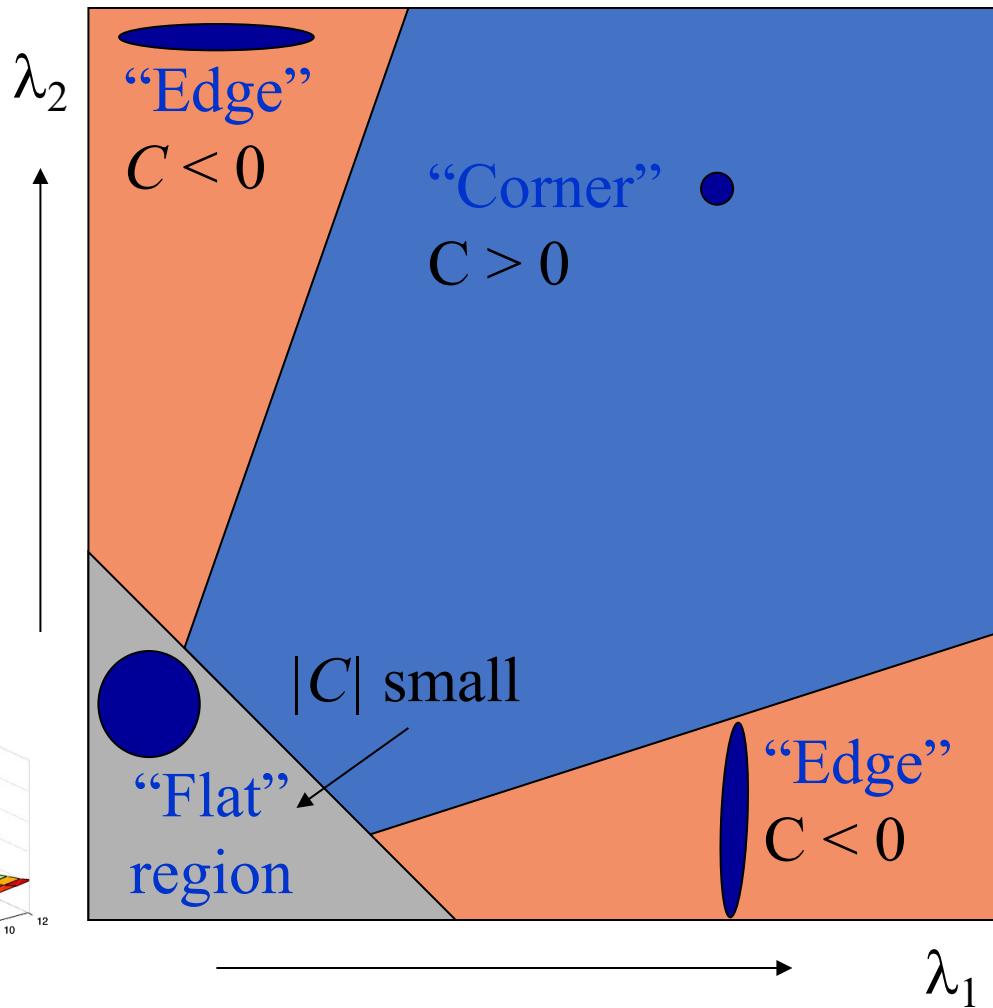
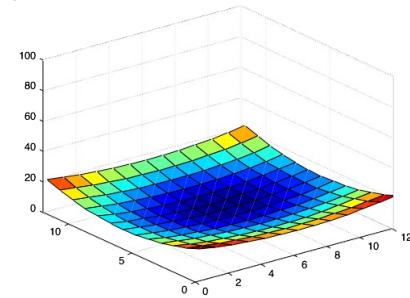


# Classification of image points using eigenvalues of $M$

Cornerness score:

$$C = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

$\alpha$ : some constant ( $\sim 0.04$  to  $0.06$ )



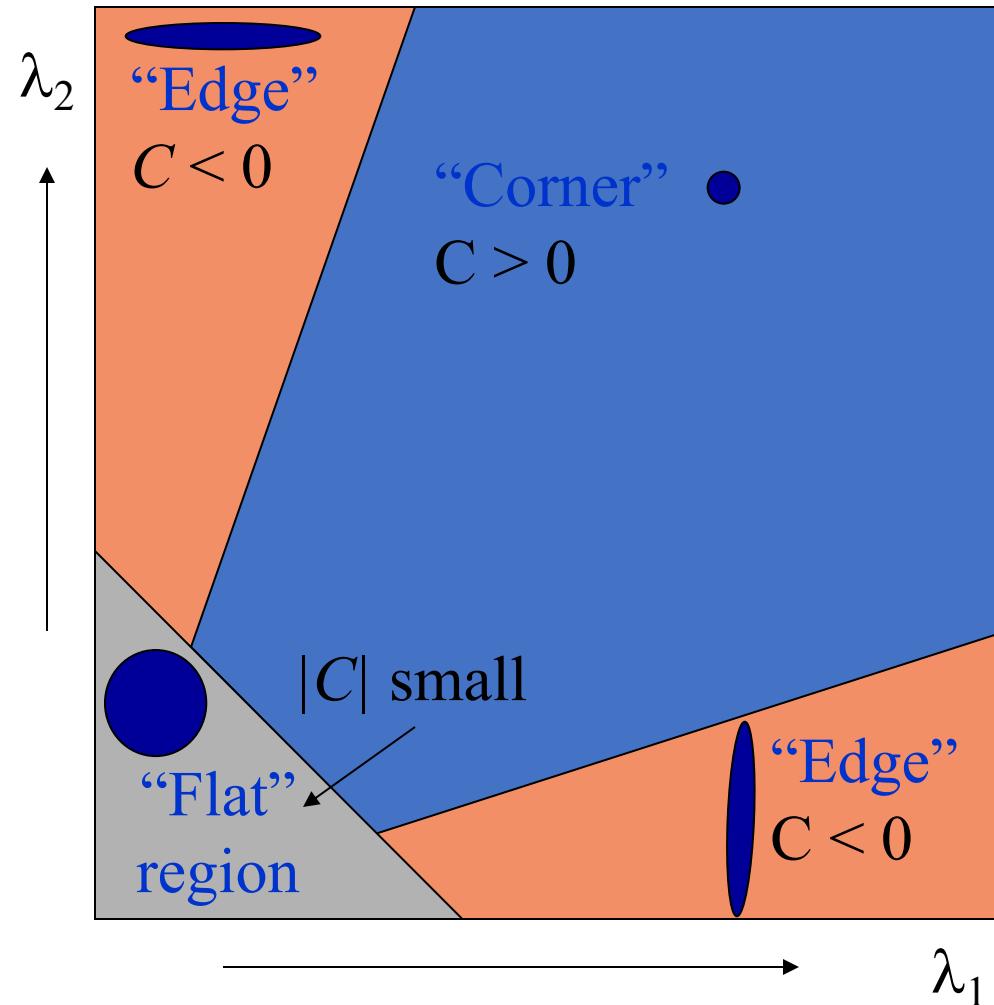
5. Use threshold on eigenvalues to detect corners

# Classification of image points using eigenvalues of $M$

Cornerness score:

$$C = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

$\alpha$ : some constant ( $\sim 0.04$  to  $0.06$ )



# Linear Algebra: Recap

Determinant of a diagonal matrix is the *product* of all eigenvalues.

$A$  is diagonal.

Trace of a square matrix is the *sum* of its diagonal entries; and is the sum of its eigenvalues.

# Classification of image points using eigenvalues of $M$

Cornerness score:

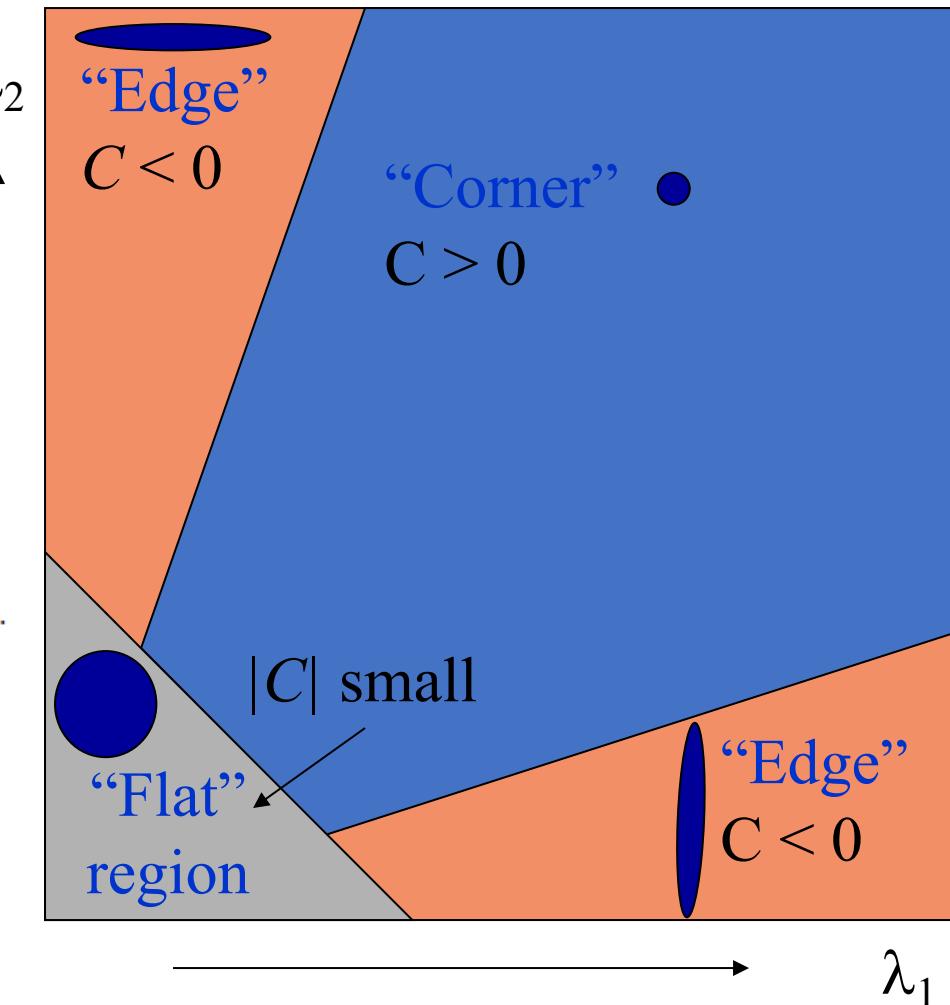
$$C = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

$\alpha$ : some constant ( $\sim 0.04$  to  $0.06$ )

Remember your linear algebra:

Determinant:  $\det(A) = \prod_{i=1}^n \lambda_i = \lambda_1 \lambda_2 \dots \lambda_n$ .  
(diagonal matrices)

Trace:  $\text{tr}(A) = \sum_i \lambda_i$ .



$$C = \det(A) - \alpha \text{trace}(A)^2$$

$$\det M = \lambda_1 \lambda_2$$

$$\text{trace } M = \lambda_1 + \lambda_2$$

$$\det \begin{pmatrix} a & b \\ c & d \end{pmatrix} = ad - bc$$

$$\text{trace} \begin{pmatrix} a & b \\ c & d \end{pmatrix} = a + d$$

# Linear Algebra: Recap

But I don't know  $A$  directly without eigen-decomposition.

- $M$  exhibits *similarity invariance* – key properties of  $M$  are invariant under similarity transforms.
- Determinant and trace of  $M$  are equal to those of  $A$ .

If  $M$  is square and diagonalizable into  $M = R^{-1}AR$ , then:

$$\det(M) = \det(R^{-1}) \det(A) \det(R)$$

$$\det(M) = \det(A) \quad \text{as} \quad \det(R^{-1}) \det(R) = 1$$

For trace, given matrices  $B, C$ , then  $\text{trace}(BC) = \text{trace}(CB)$

$$\text{trace}(R^{-1}(AR)) = \text{trace}((AR)R^{-1}) = \text{trace}(A)$$

so  $\text{trace}(M) = \text{trace}(A)$

# Classification of image points using eigenvalues of $M$

Cornerness score:

$$C = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

$\alpha$ : some constant ( $\sim 0.04$  to  $0.06$ )

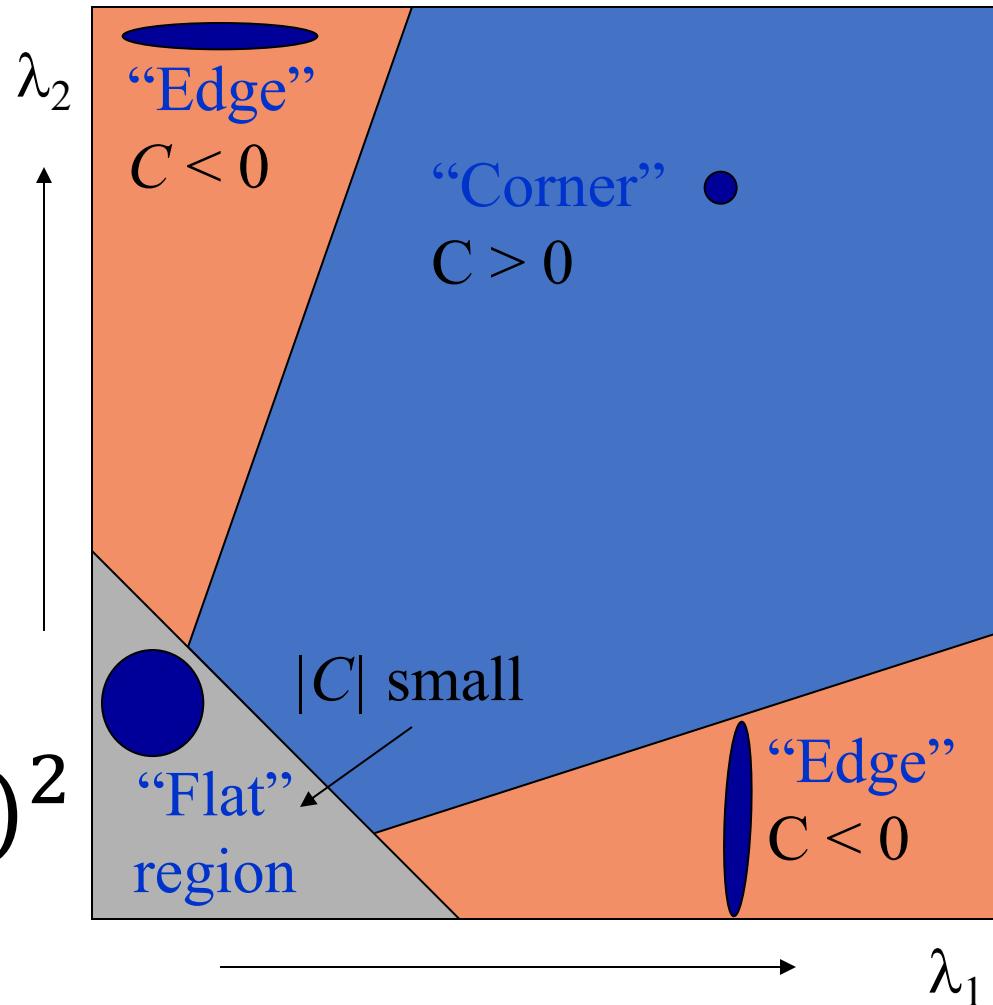
*Remember your linear algebra:*

Determinant:  $\det(A) = \prod_{i=1}^n \lambda_i = \lambda_1 \lambda_2 \dots \lambda_n$ .  
(diagonal matrices)

Trace:  $\text{tr}(A) = \sum_i \lambda_i$ .

$$C = \det(M) - \alpha \text{trace}(M)^2$$

Avoids explicit eigenvalue computation!



Harris & Stephens (1988)

$$R = \det(M) - \kappa \text{trace}^2(M)$$

Kanade & Tomasi (1994)

$$R = \min(\lambda_1, \lambda_2)$$

Nobel (1998)

$$R = \frac{\det(M)}{\text{trace}(M) + \epsilon}$$

# Harris Corner Detector

1) Compute  $M$  matrix for each pixel to recover *cornerness* score  $C$ .

Note: We can find  $M$  purely from the per-pixel image derivatives!

2) Threshold to find pixels which give large corner response  
( $C > \text{threshold}$ ).

3) Find the local maxima pixels, i.e., non-maxima suppression.

# Harris Corner Detector

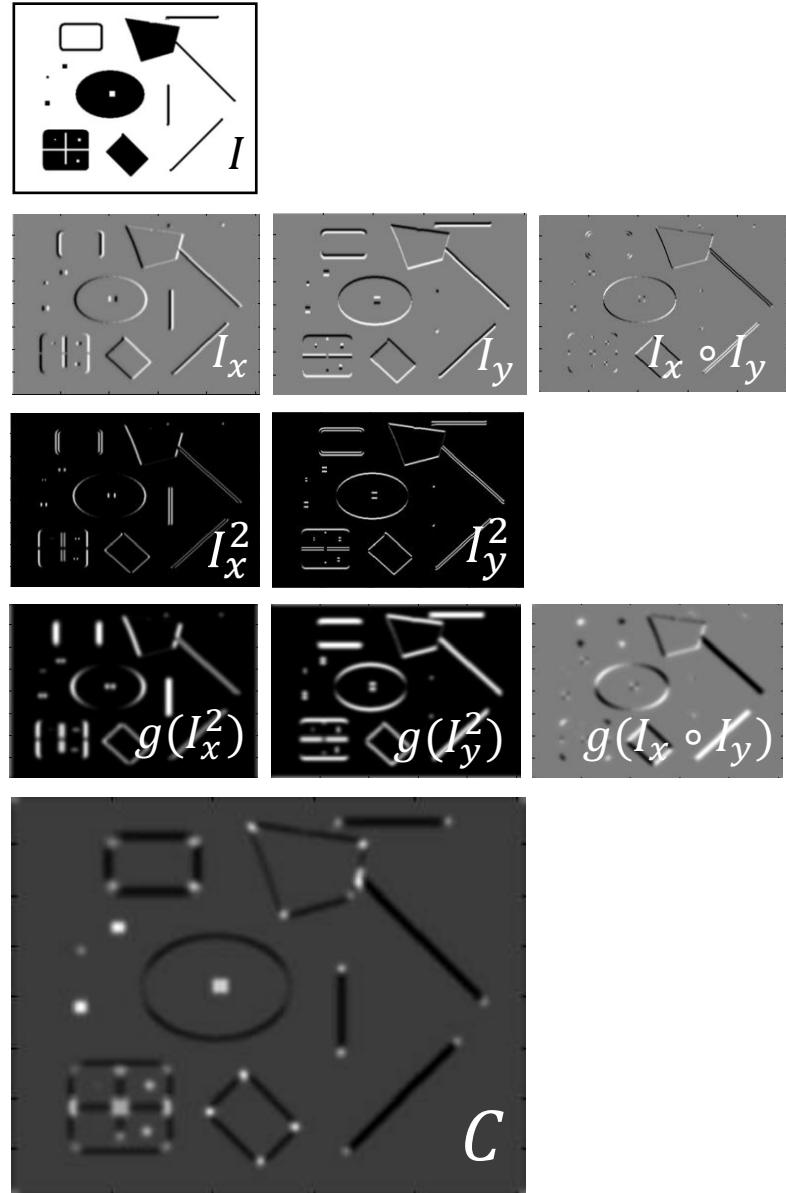
- Compute image discrete first derivatives.
- Compute products of derivatives to use in  $M$
- Create  $M$  via window function: Gaussian filter  $g()$  with  $\sigma$   

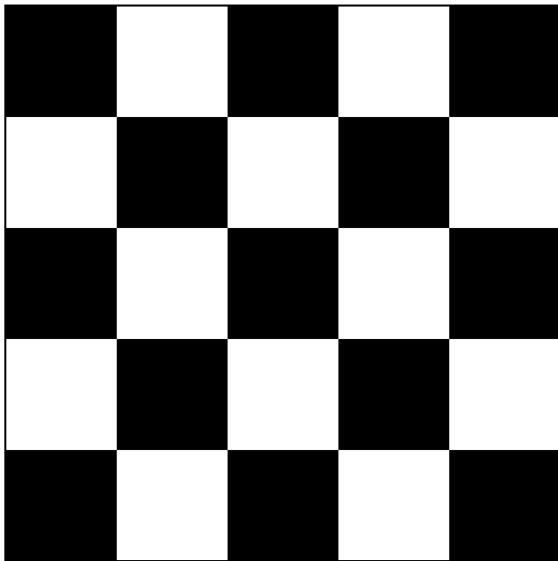
$$g(I_x^2), g(I_y^2), g(I_x \circ I_y)$$
- Compute cornerness

$$\begin{aligned} C = \det(M) - \alpha \operatorname{trace}(M)^2 &= g(I_x^2) \circ g(I_y^2) - g(I_x \circ I_y)^2 \\ &\quad - \alpha [g(I_x^2) + g(I_y^2)]^2 \end{aligned}$$

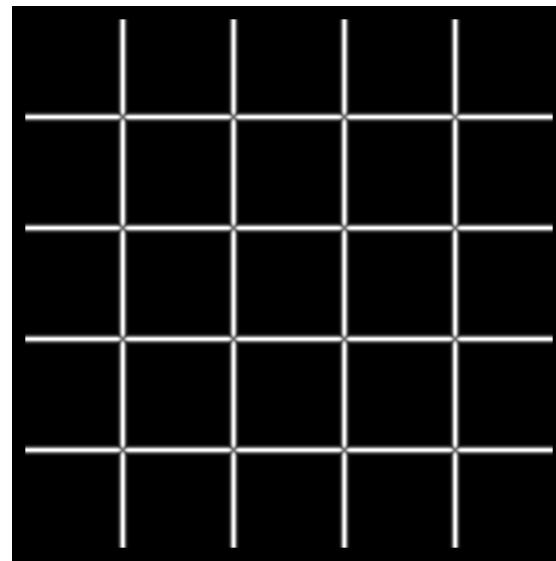
Reminder:  $a \circ b$  is  
Hadamard product  
(element-wise  
multiplication)

5. Threshold on  $C$  to pick high cornerness
6. Non-maximal suppression to pick peaks.

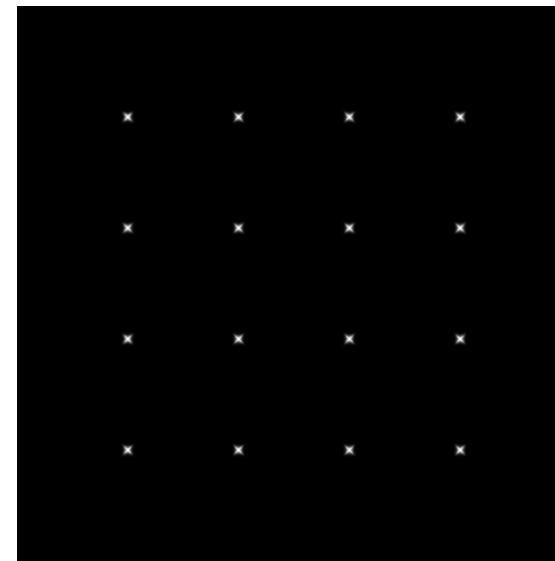




$I$



$\lambda_{\max}$



$\lambda_{\min}$

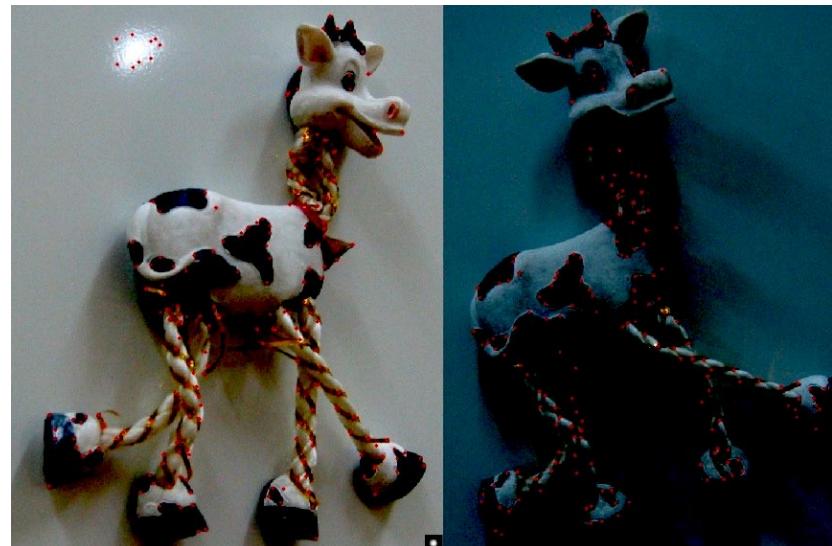
Yet another option:

$$f = \frac{\lambda_1 \lambda_2}{\lambda_1 + \lambda_2} = \frac{\text{determinant}(H)}{\text{trace}(H)}$$

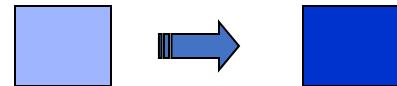
# Invariance and Covariance

Are locations *invariant* to photometric transformations and *covariant* to geometric transformations?

- **Invariance:** image is transformed and corner locations do not change
- **Covariance / Equivariance:** if we have two transformed versions of the same image, features should be detected in corresponding locations



# Affine Intensity Change



$$I \rightarrow a I + b$$

Only derivatives are used => invariance to intensity shift  $I \rightarrow I + b$

Intensity scaling:  $I \rightarrow a I \Rightarrow$  still affected

Image 1 row

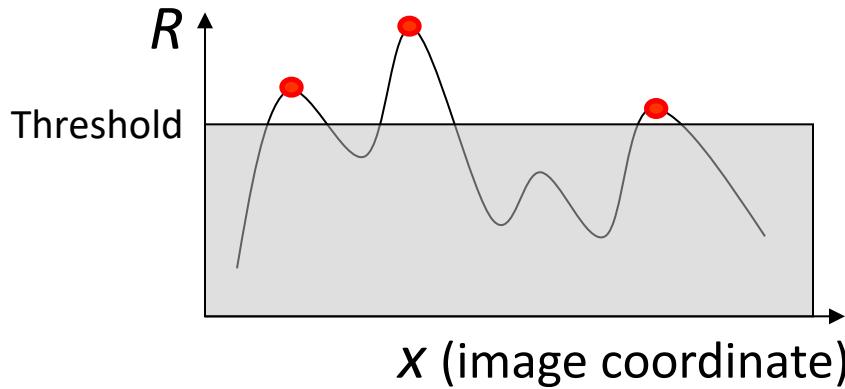
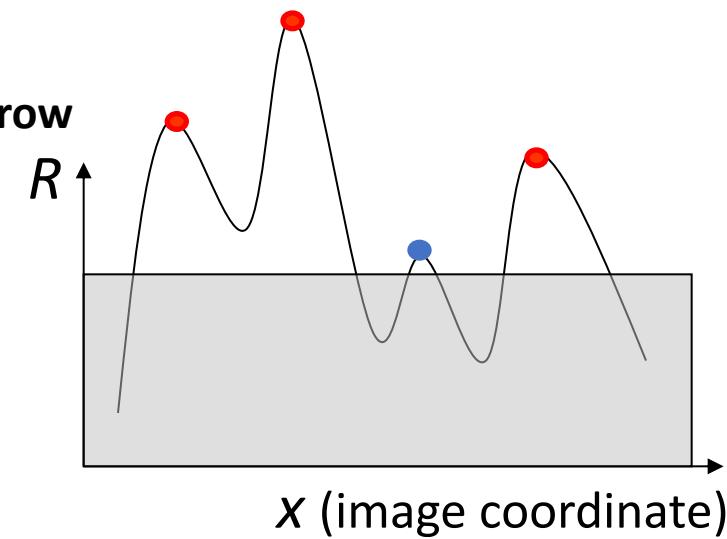
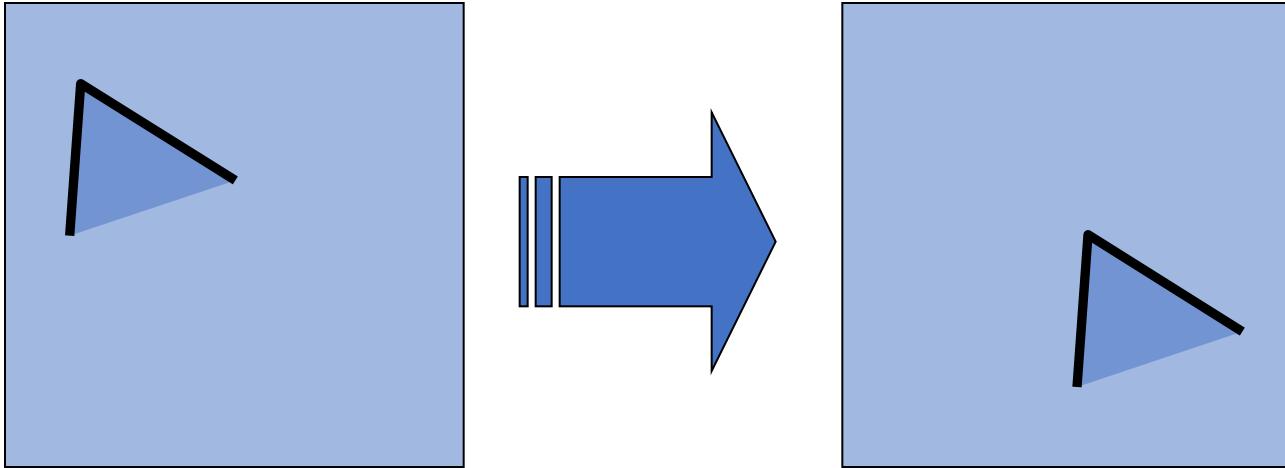


Image 2 row



*Partially invariant to affine intensity change*

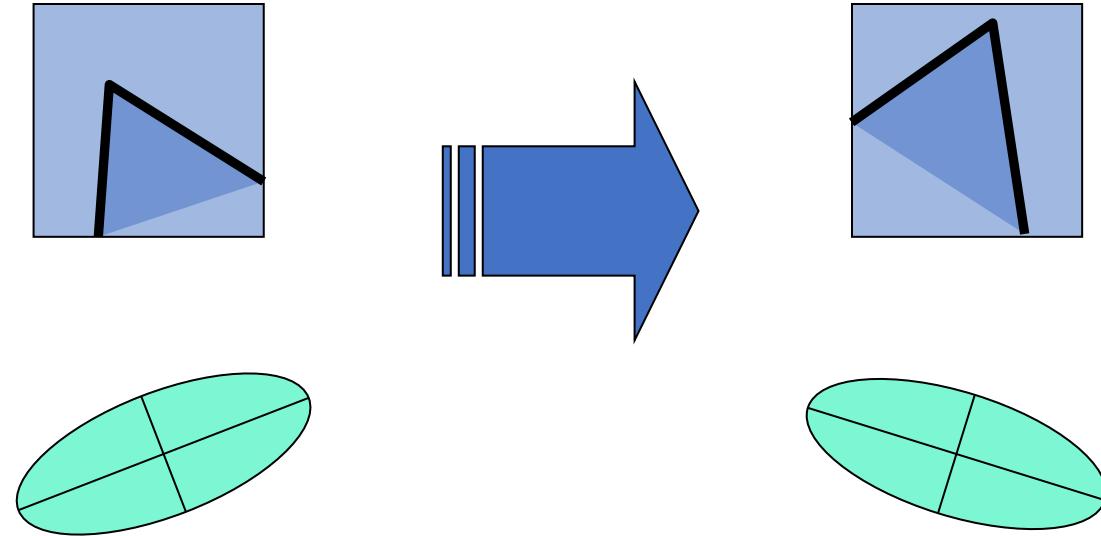
# Image Translation



- Derivatives and window function are shift-invariant.

Corner location is covariant w.r.t. translation

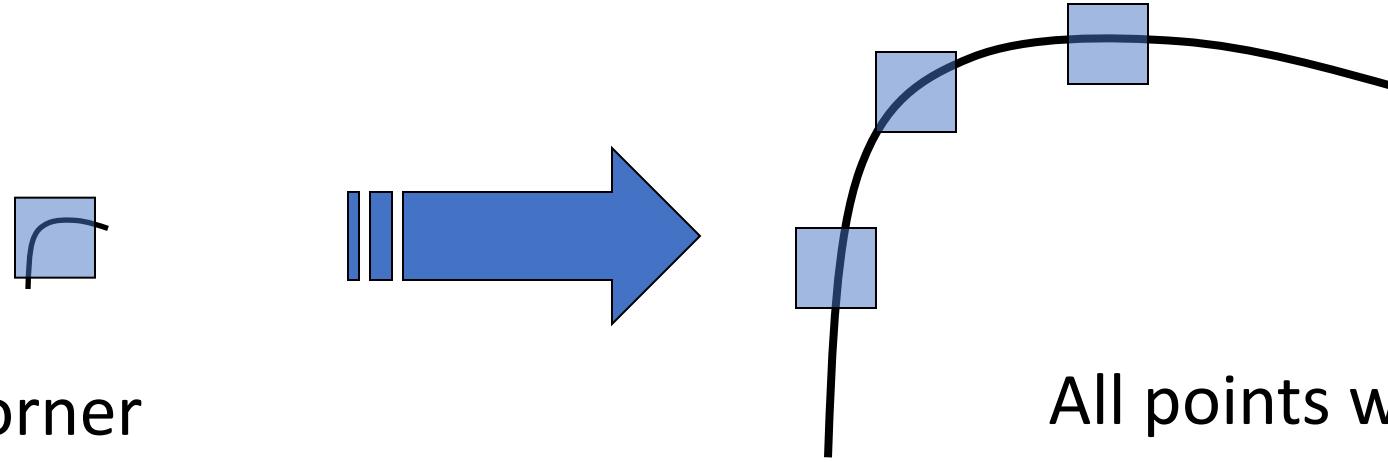
# Image Rotation



Second moment ellipse rotates but its shape (i.e., eigenvalues) remains the same.

Corner location is covariant w.r.t. rotation

# Scaling



Corner

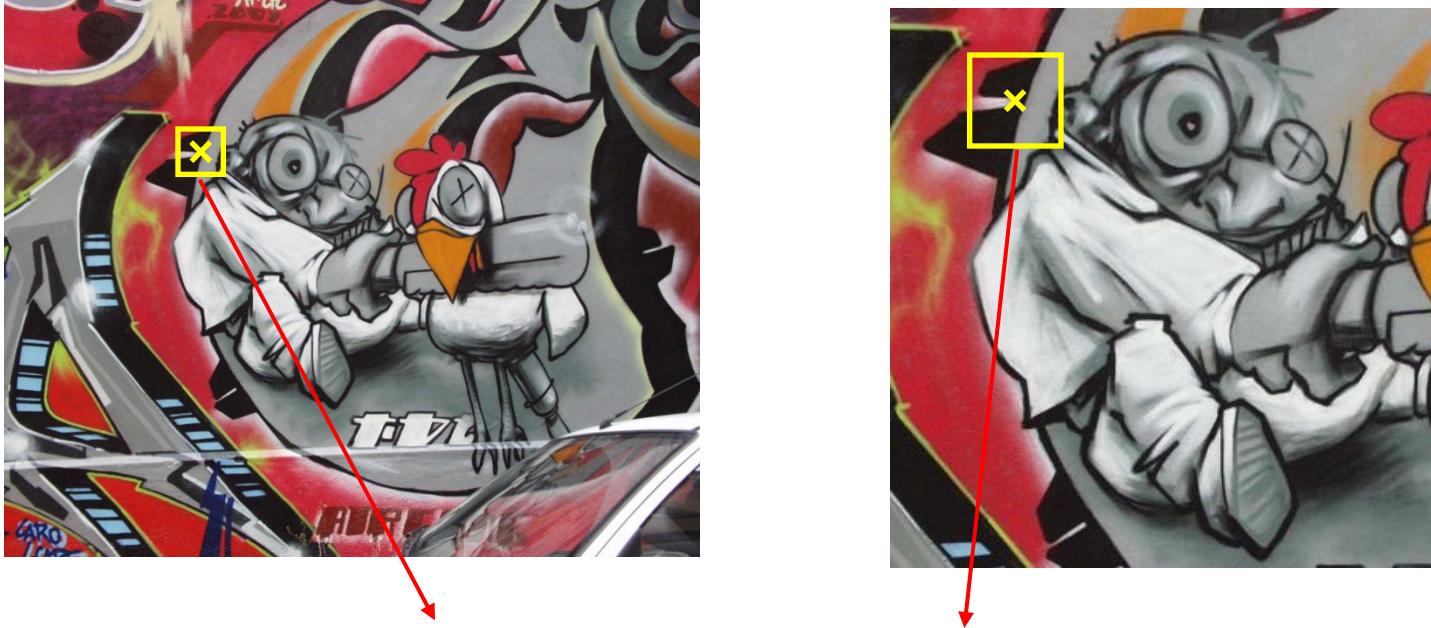
All points will be  
incorrectly  
classified as  
edges

Corner location is not covariant to scaling!

How can we make a feature detector scale-invariant?

Multi-scale detection

# Automatic Scale Selection

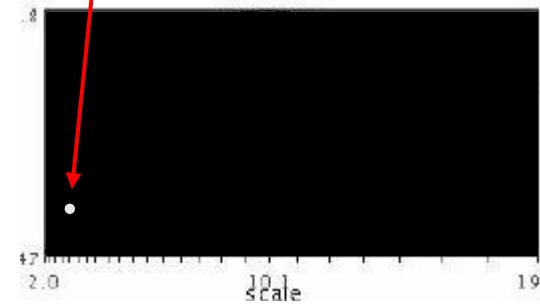
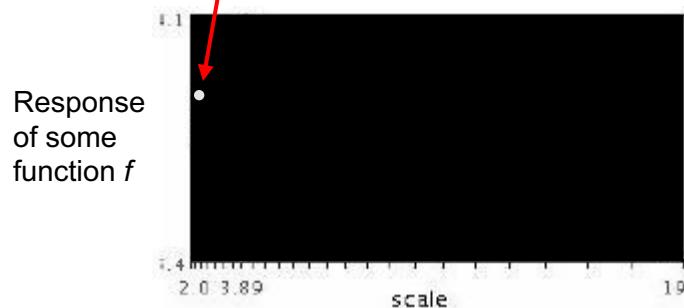


How to find patch sizes at which  $f$  response is equal?

What is a good  $f$  ?

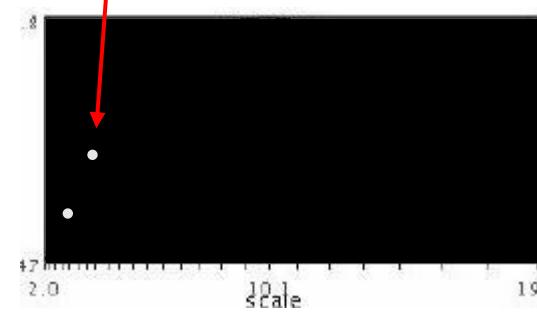
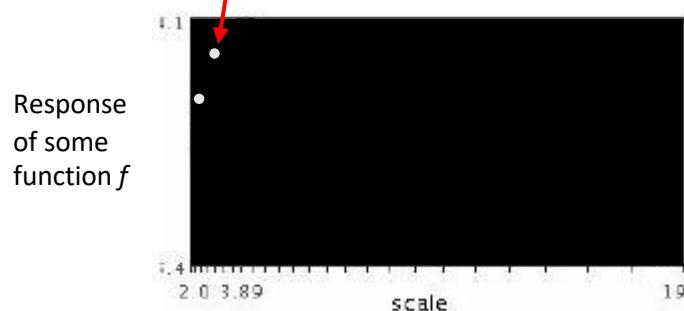
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



# Automatic Scale Selection

- Function responses for increasing scale (scale signature)

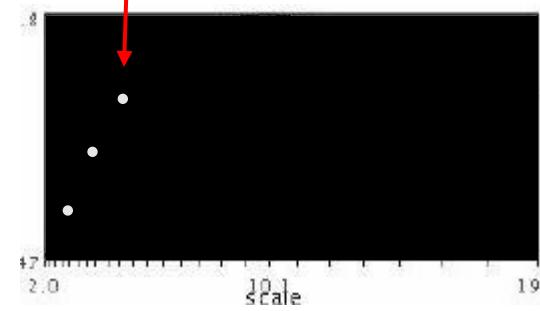
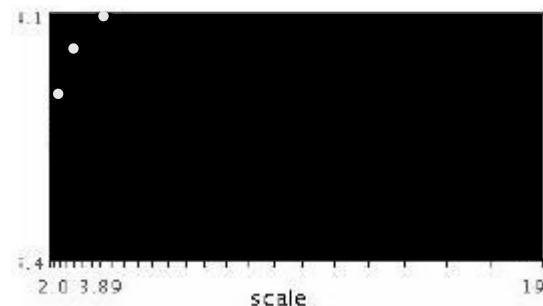


# Automatic Scale Selection

- Function responses for increasing scale (scale signature)

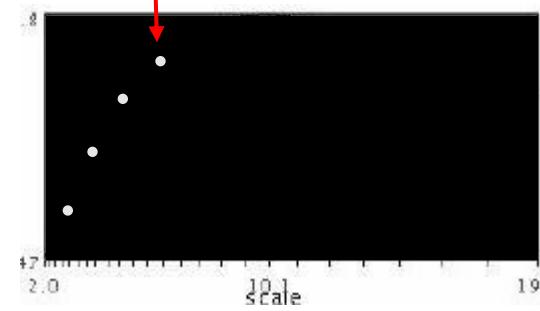
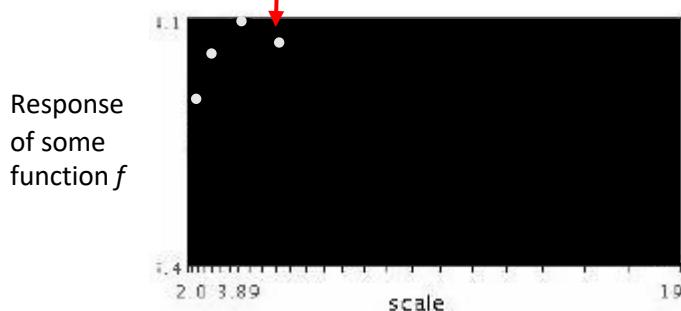
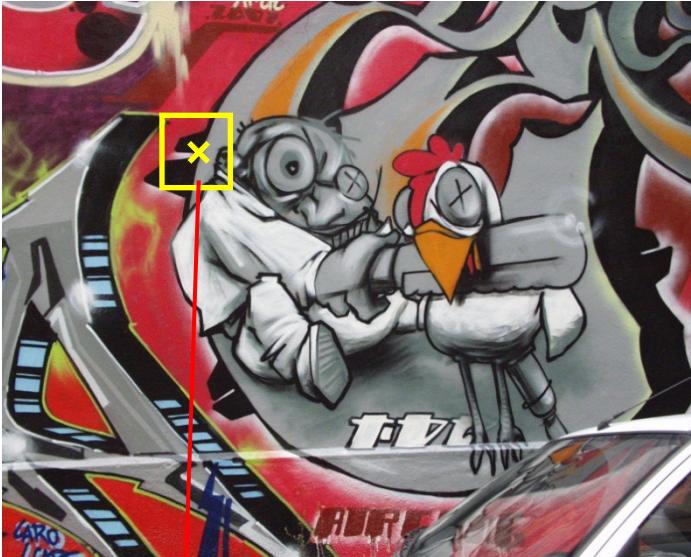


Response  
of some  
function  $f$



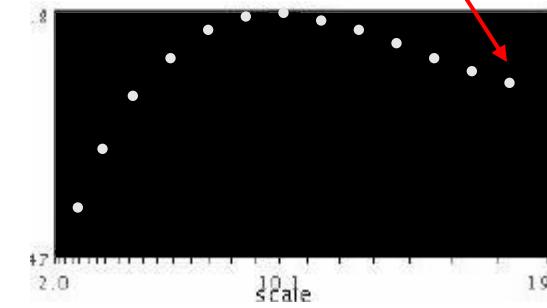
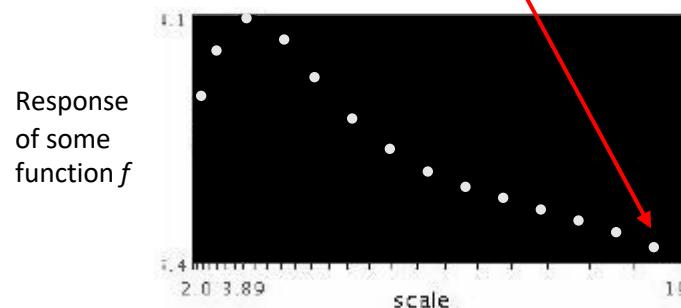
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



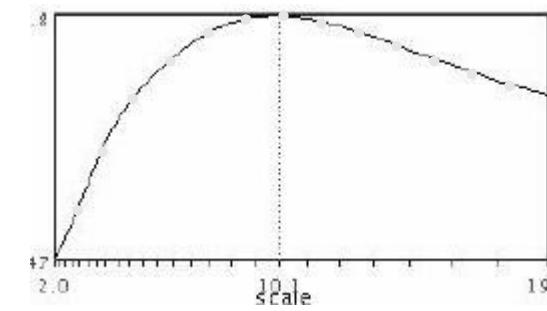
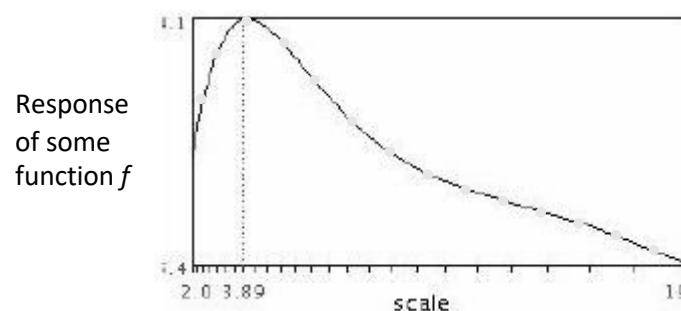
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



Intuitively...

Find local maxima in both **position** and **scale**

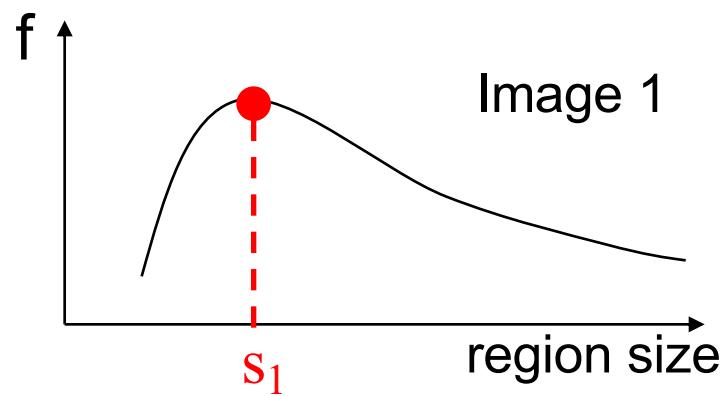
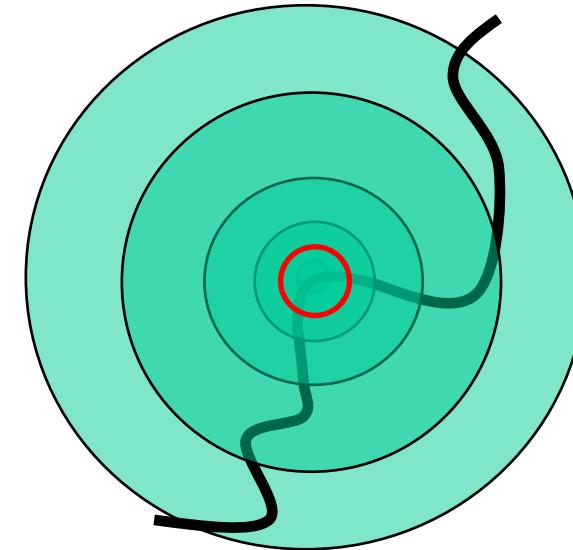
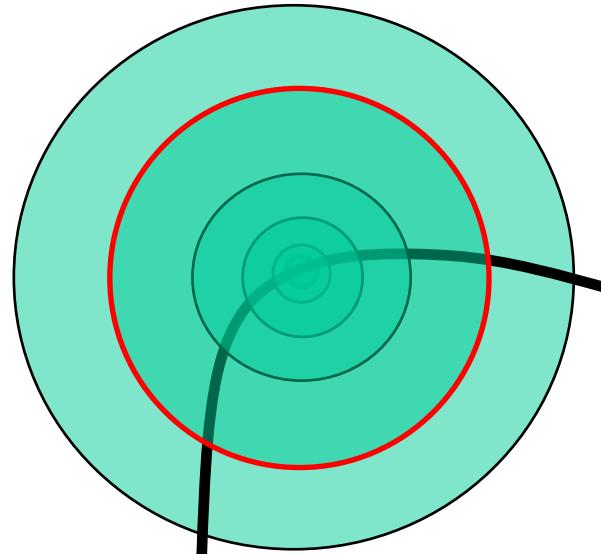


Image 1

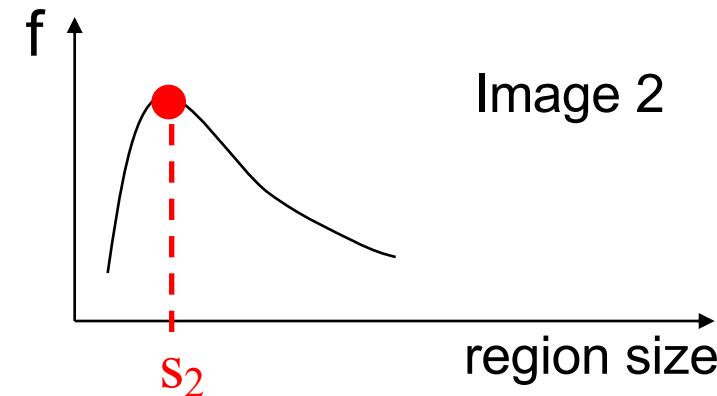
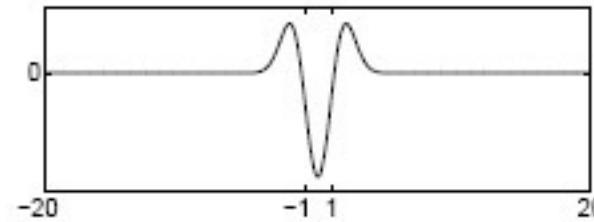


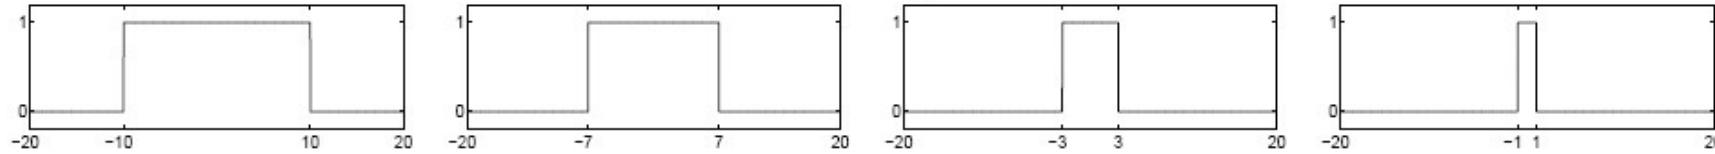
Image 2

Formally...

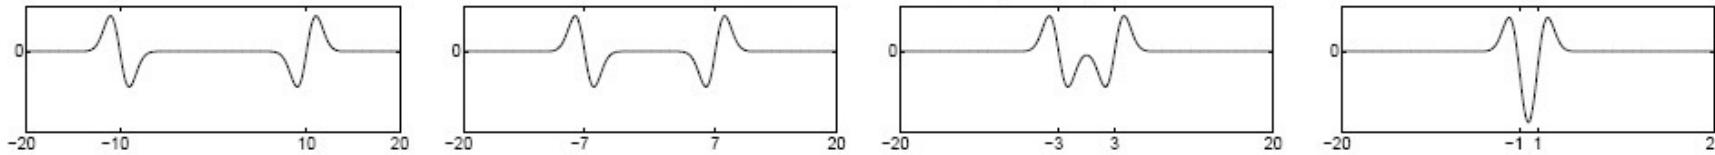
Laplacian filter



Original signal

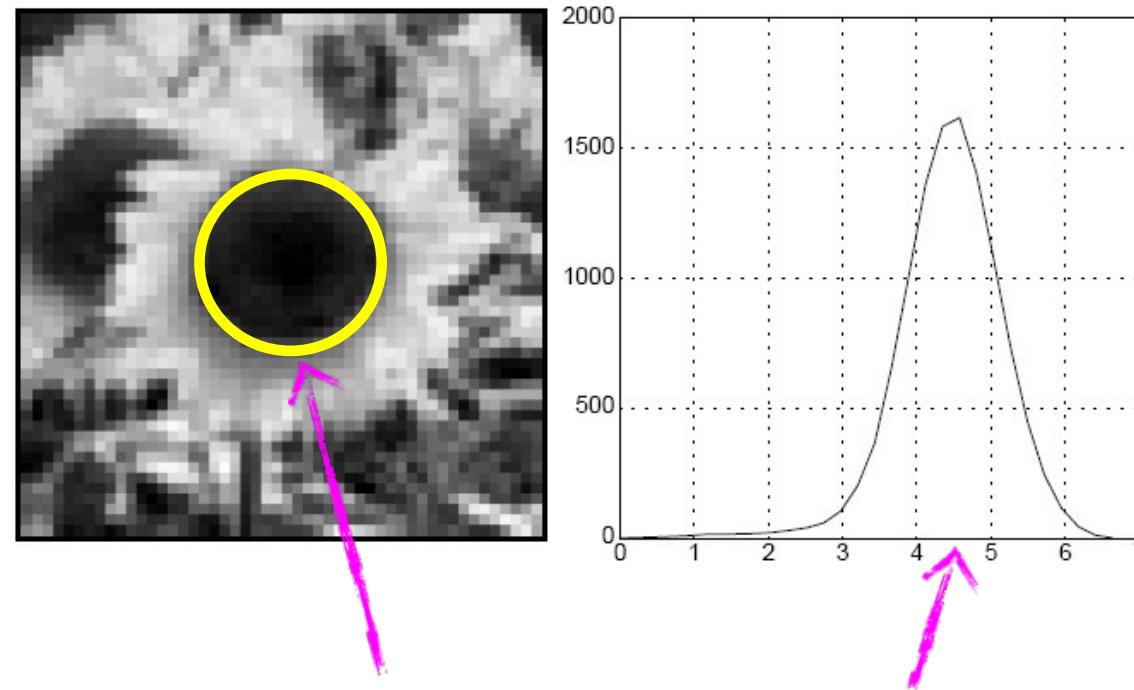


Convolved with Laplacian ( $\sigma = 1$ )



Highest response when the signal has the same **characteristic scale** as the filter

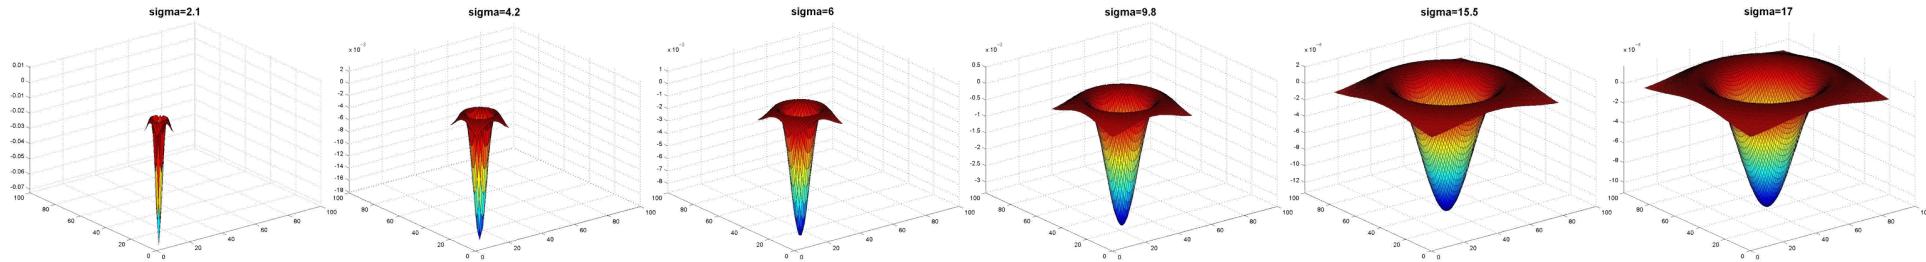
characteristic scale - the scale that produces peak filter response



characteristic scale

**we need to search over characteristic scales**

# What happens if you apply different Laplacian filters?



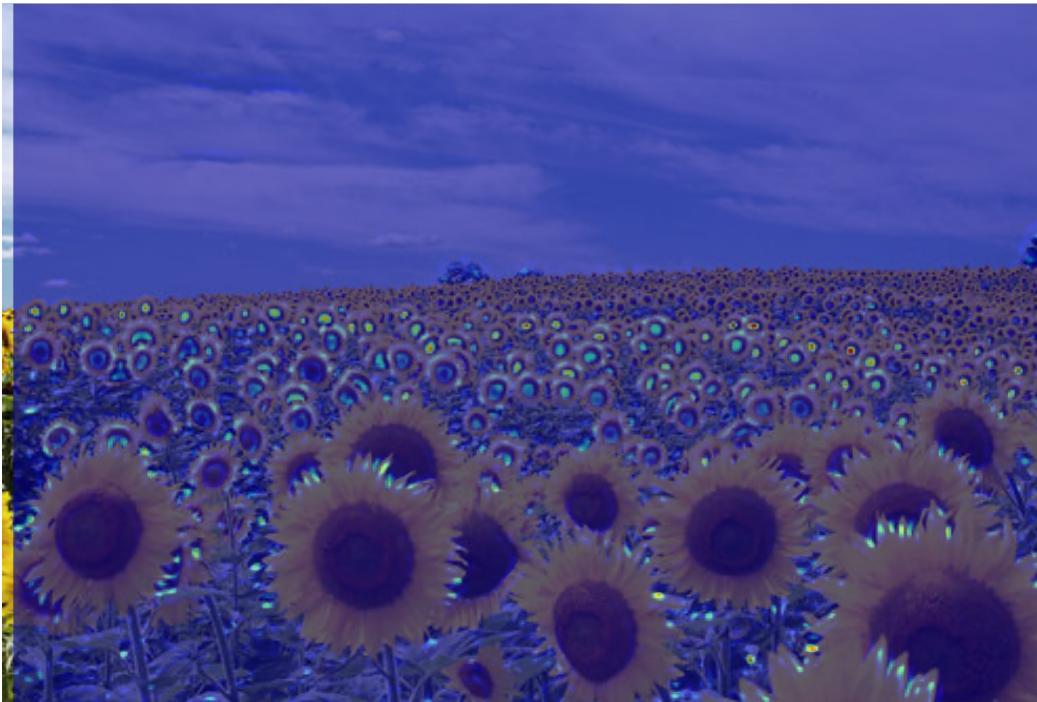
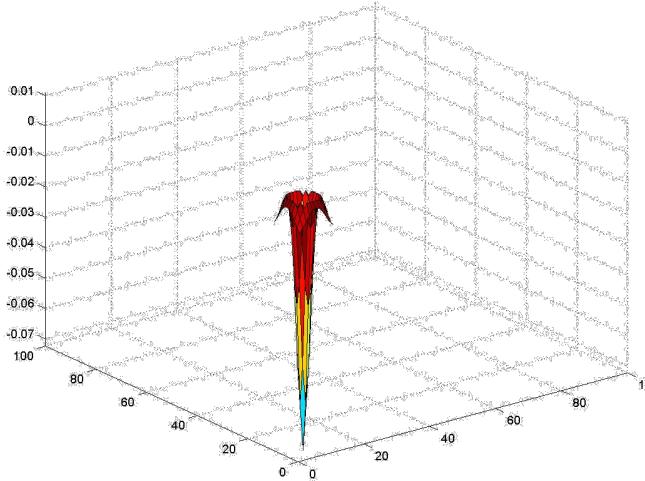
Full size



3/4 size

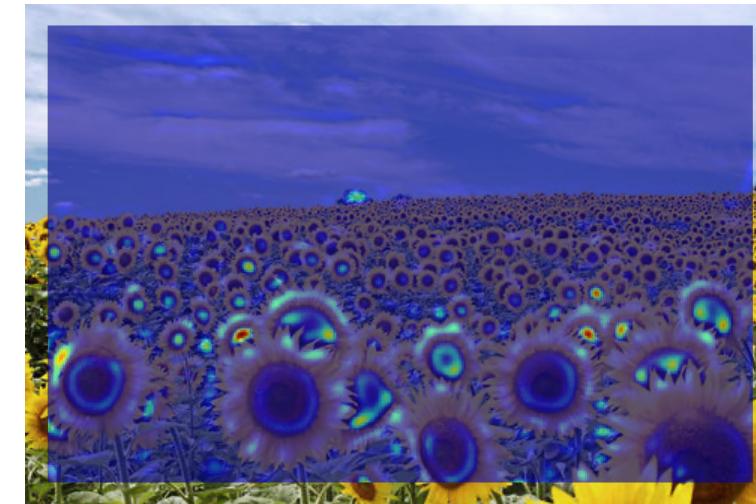
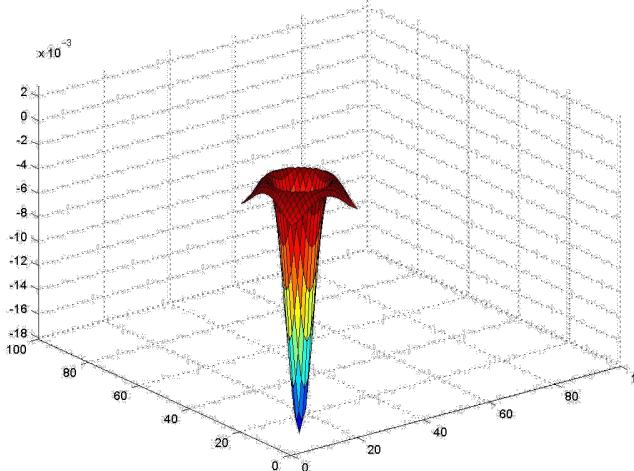


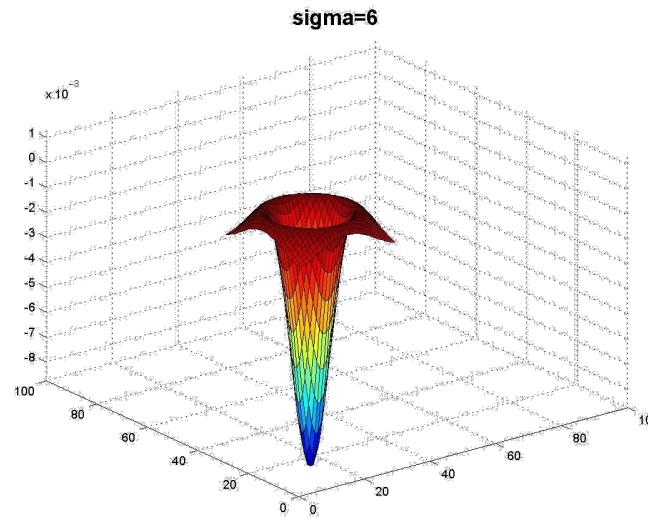
**sigma=2.1**



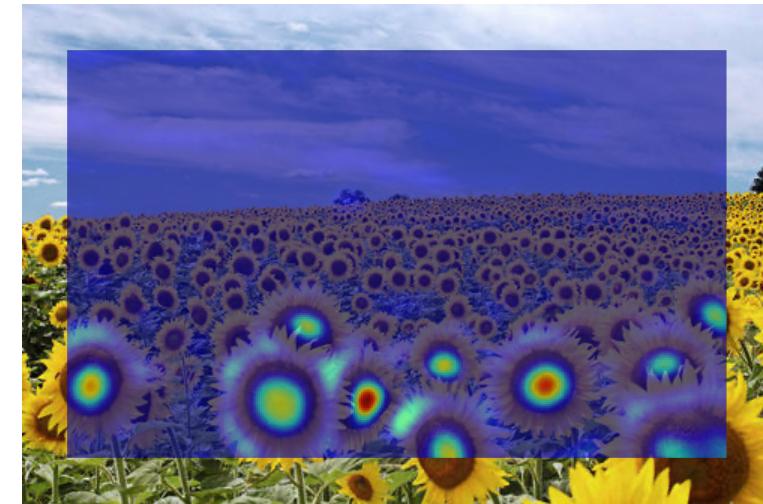
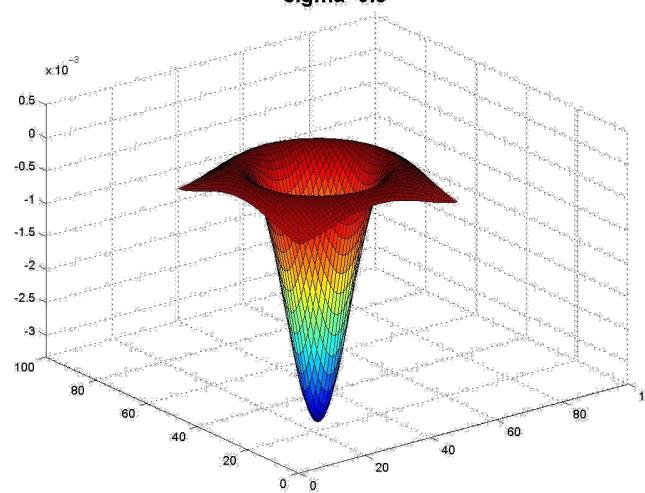
**jet** color scale  
blue: low, red: high

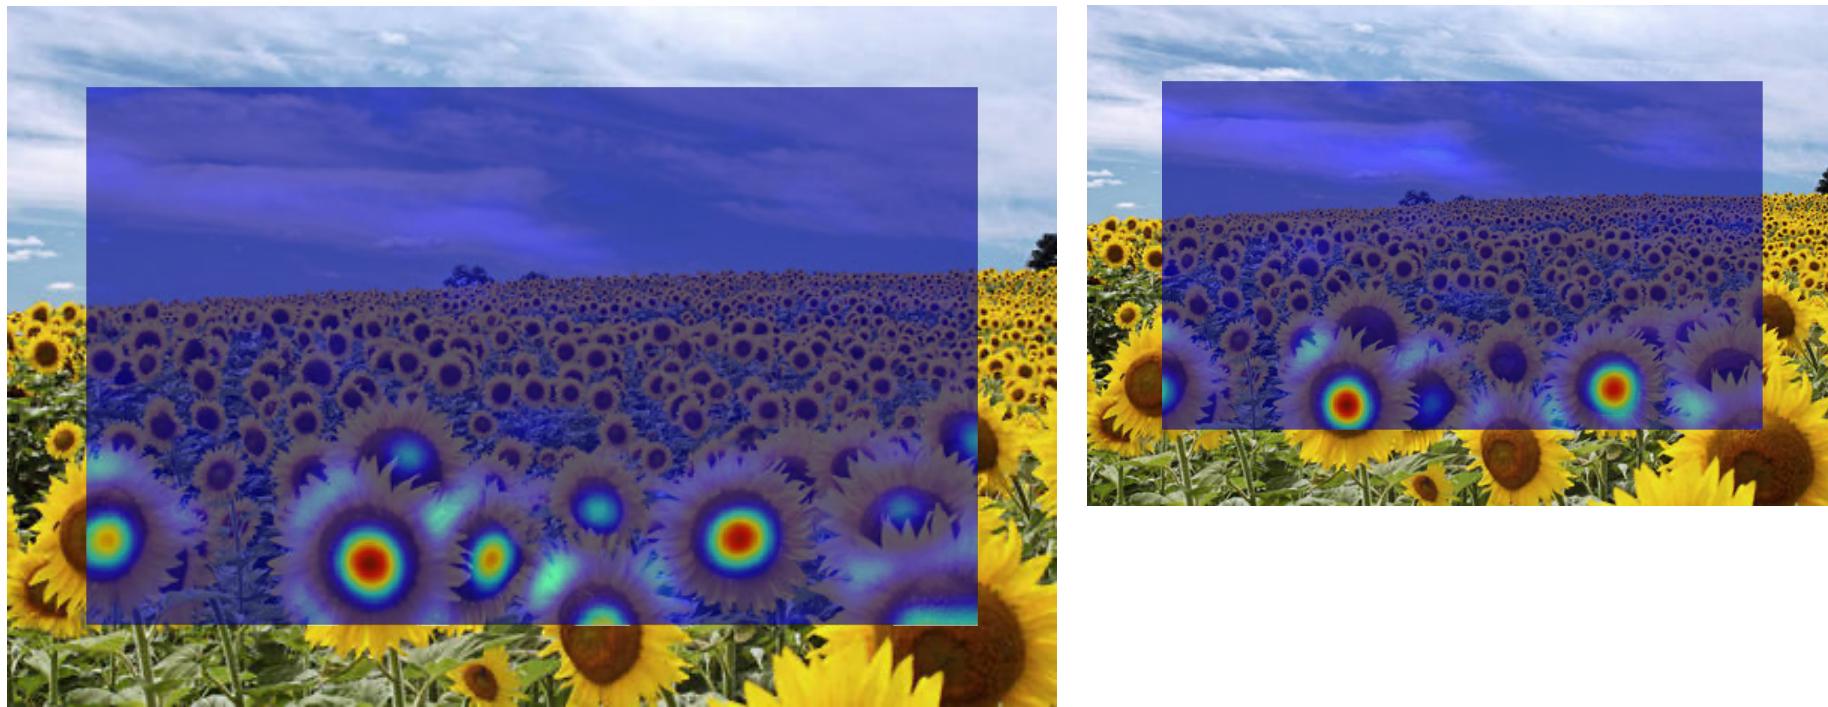
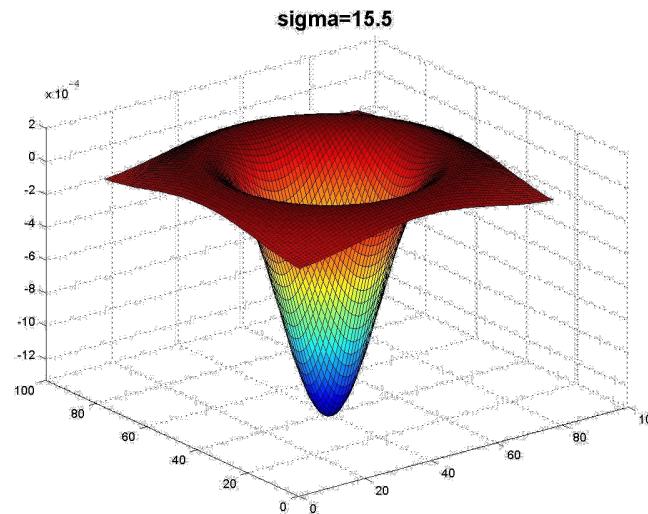
**sigma=4.2**



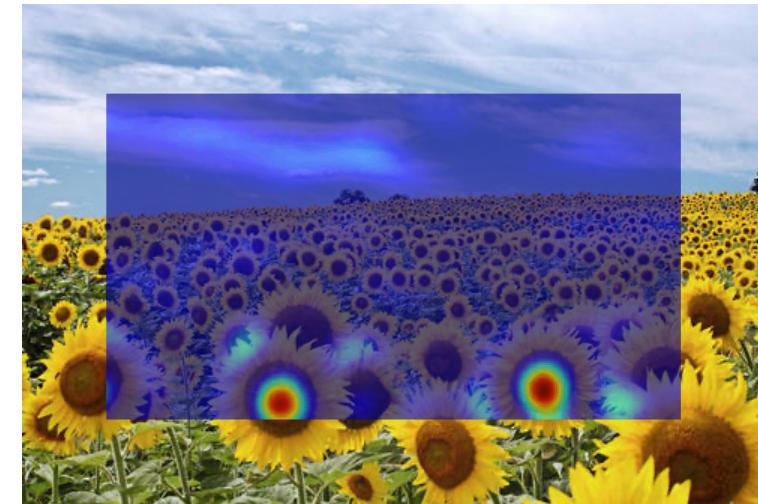
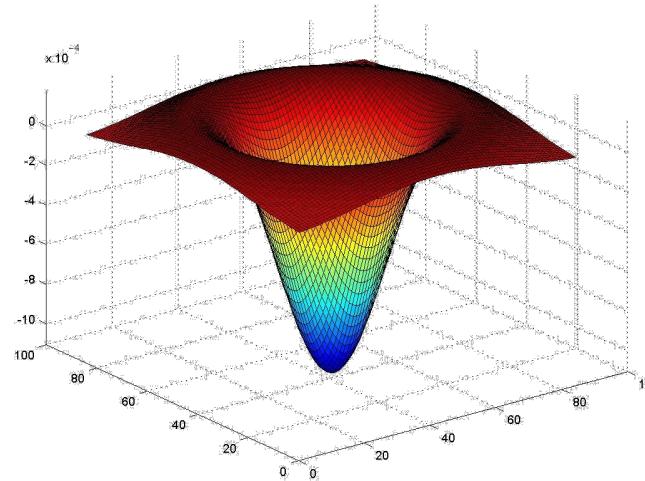


**sigma=9.8**





**sigma=17**



What happened when you applied different Laplacian filters?

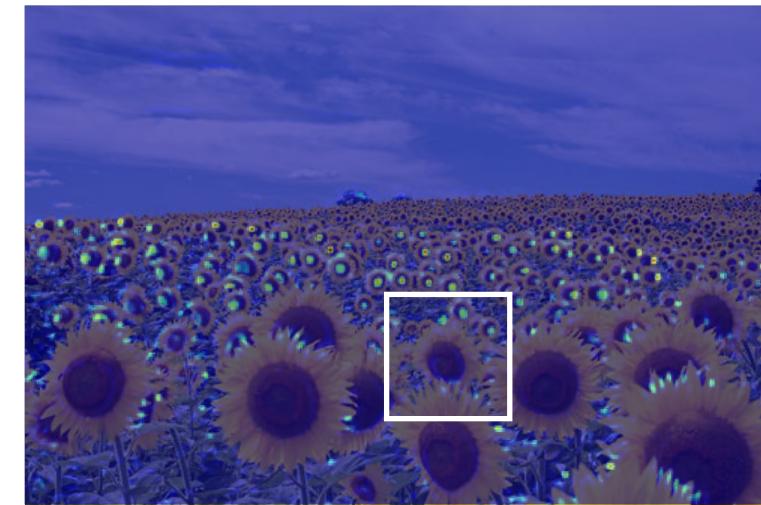
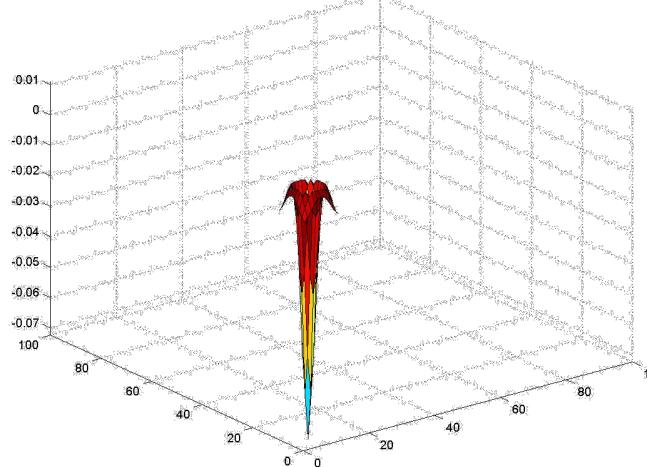
Full size



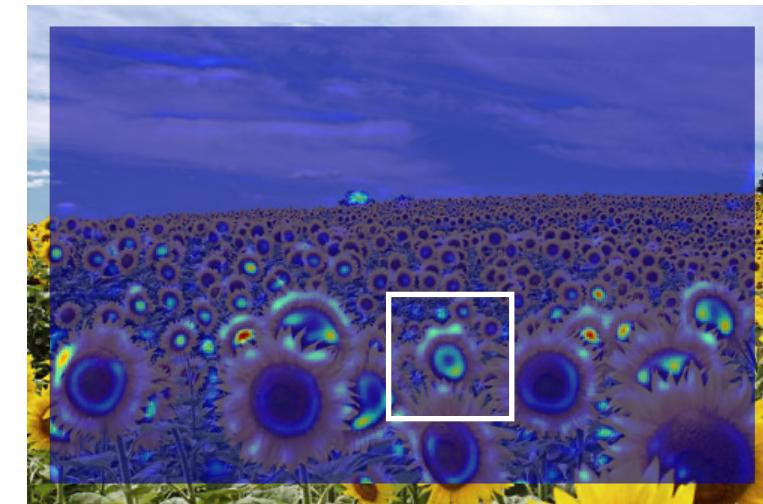
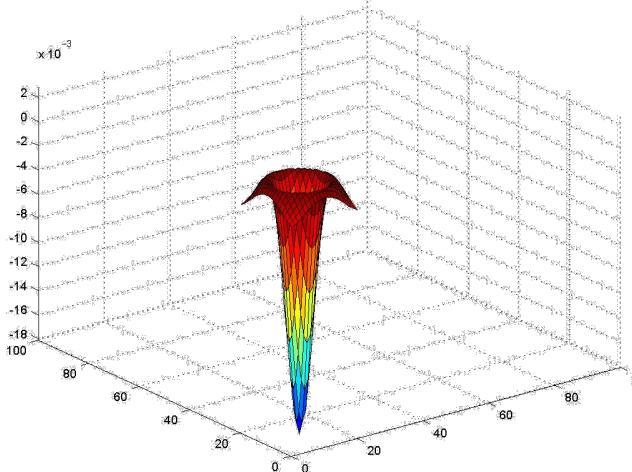
3/4 size

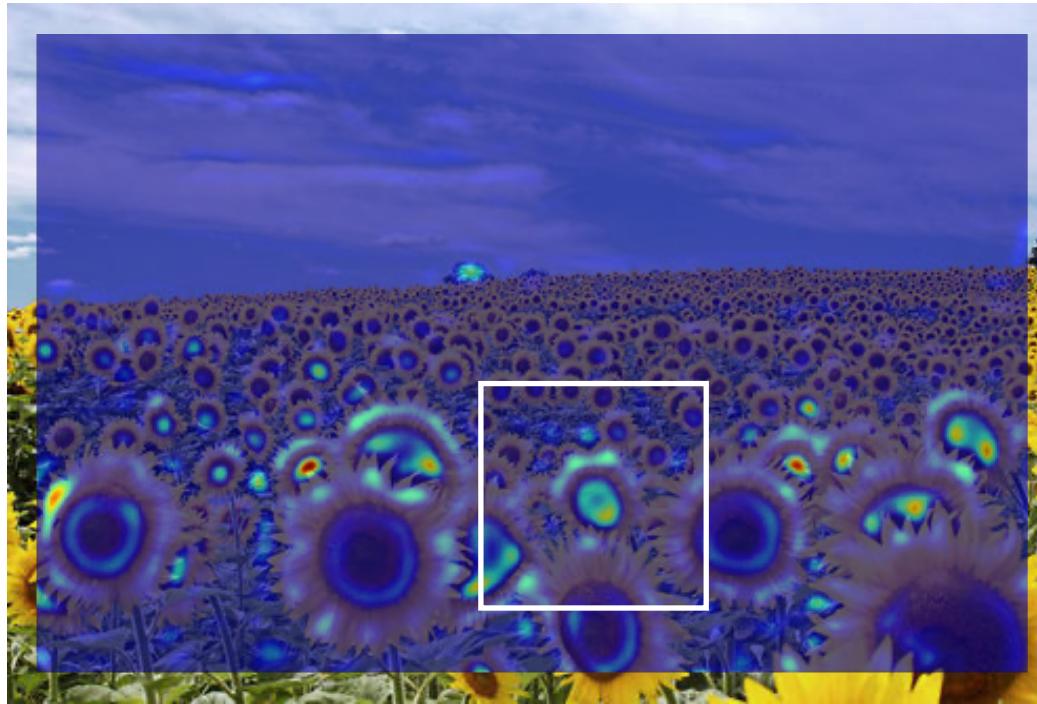
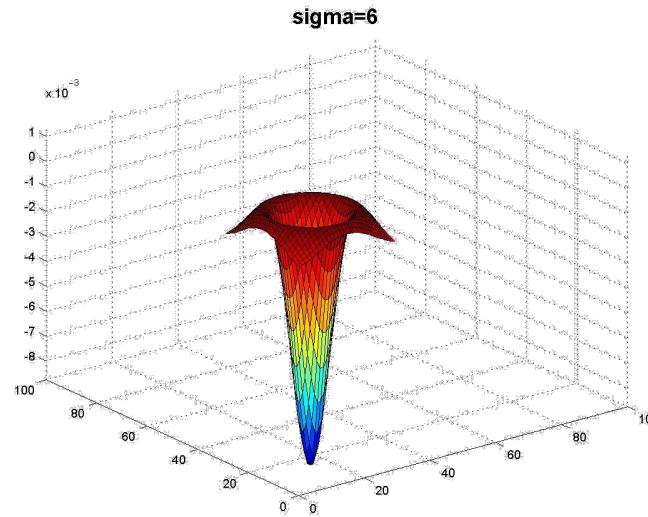


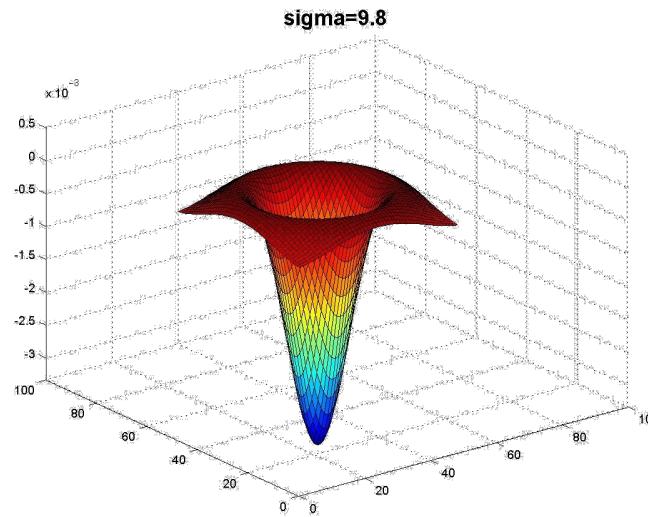
**sigma=2.1**

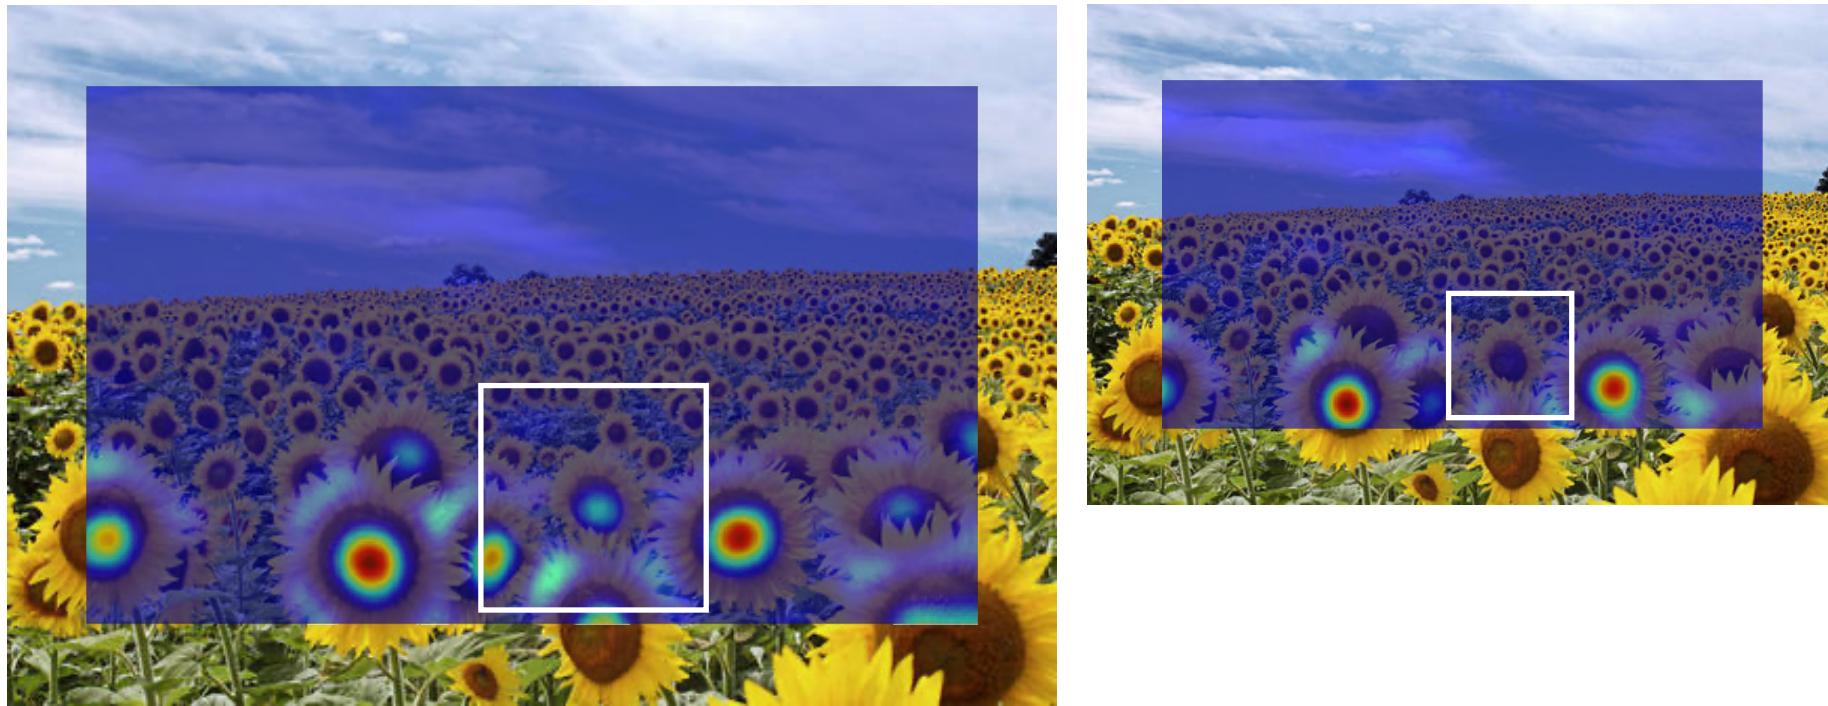
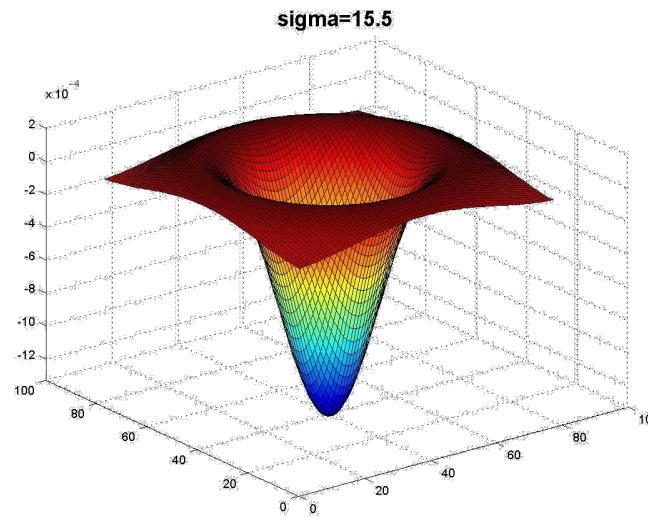


**sigma=4.2**

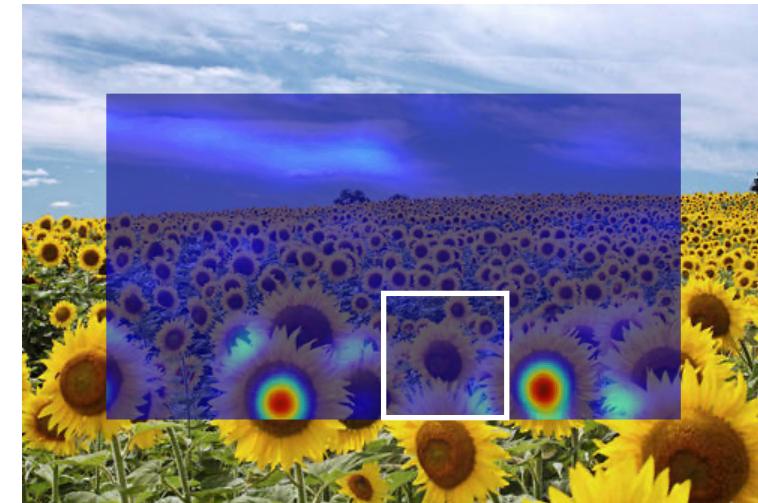
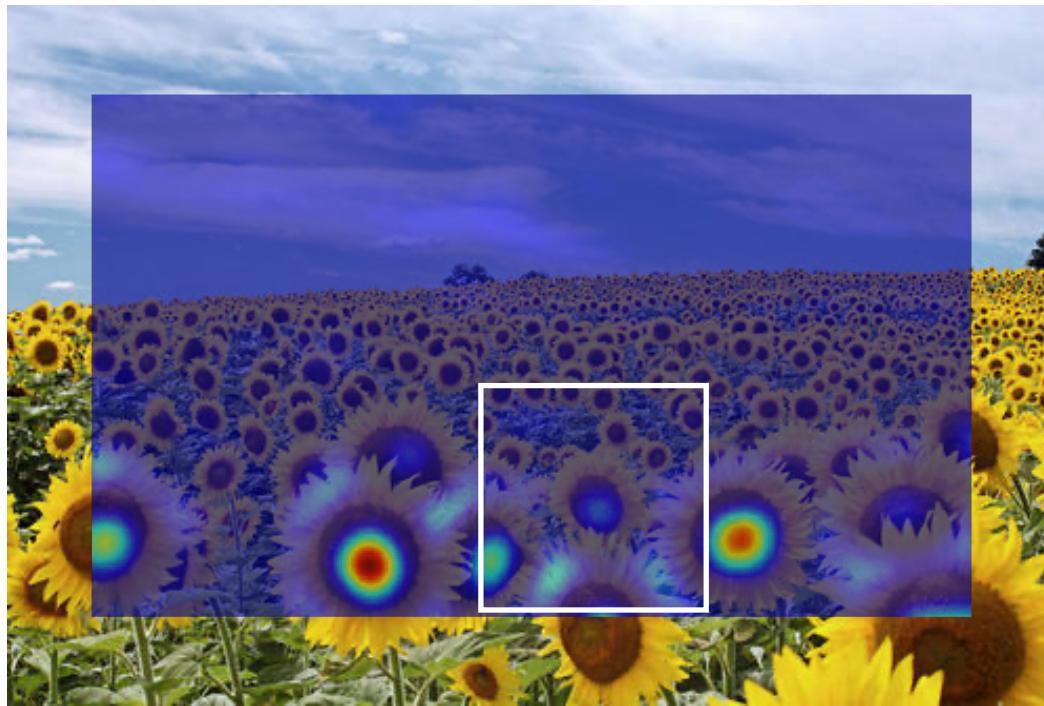
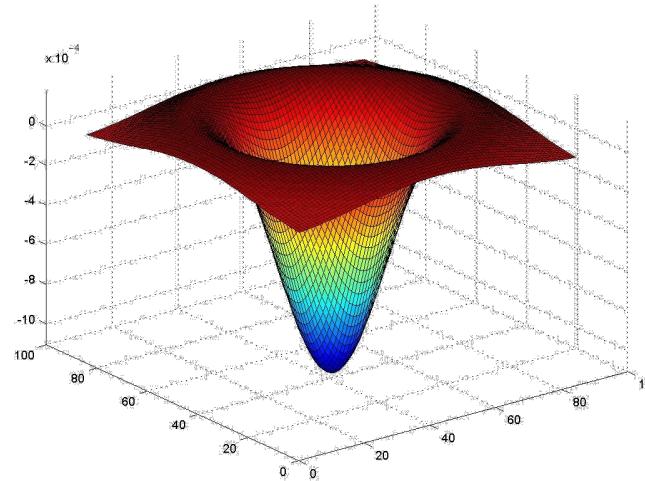








**sigma=17**



What happened when you applied different Laplacian filters?

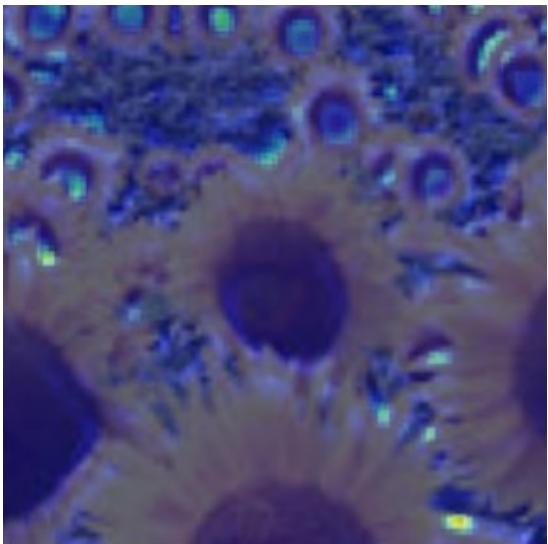
Full size



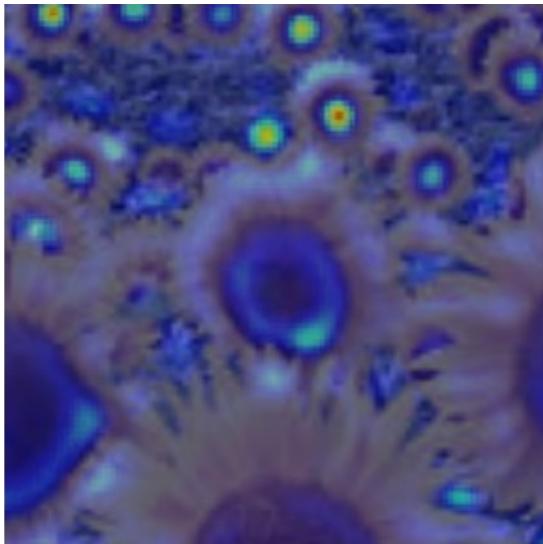
3/4 size



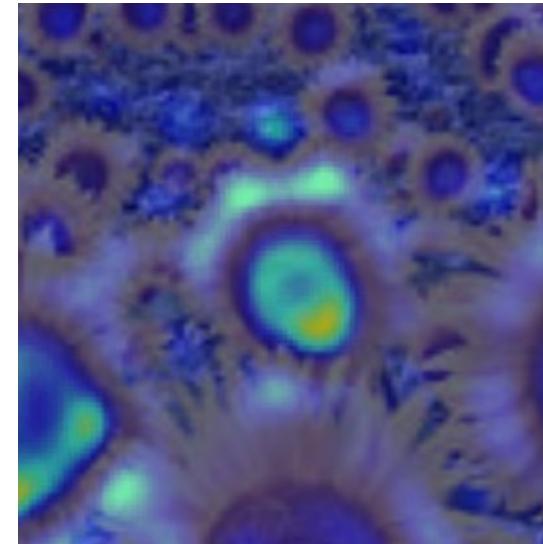
2.1



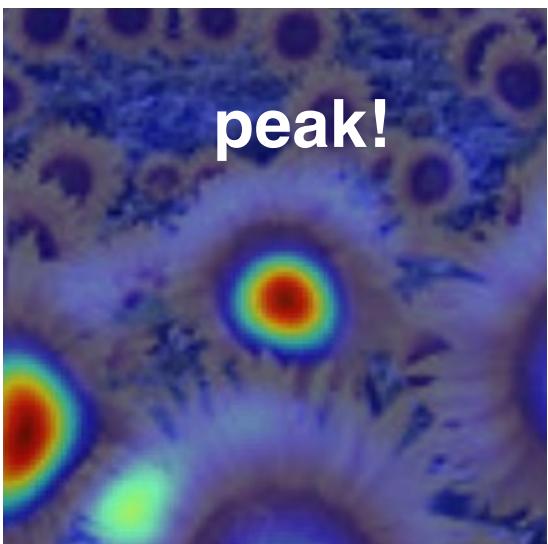
4.2



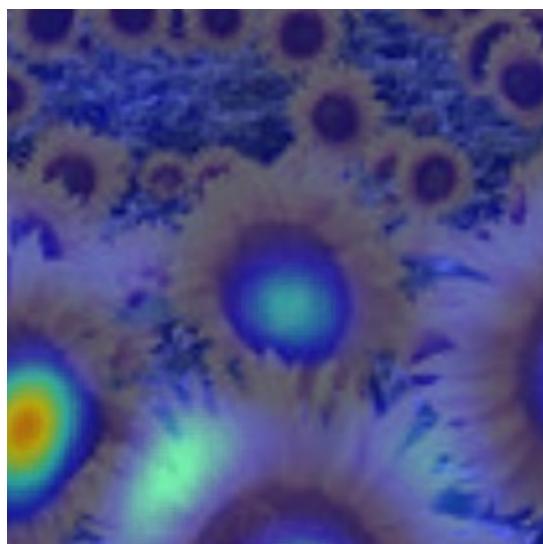
6.0



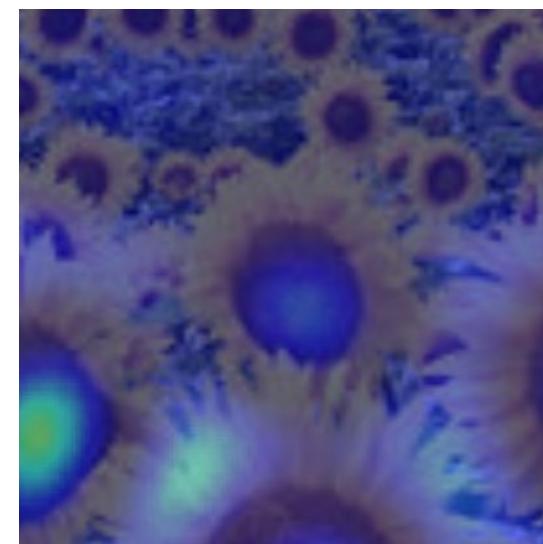
9.8



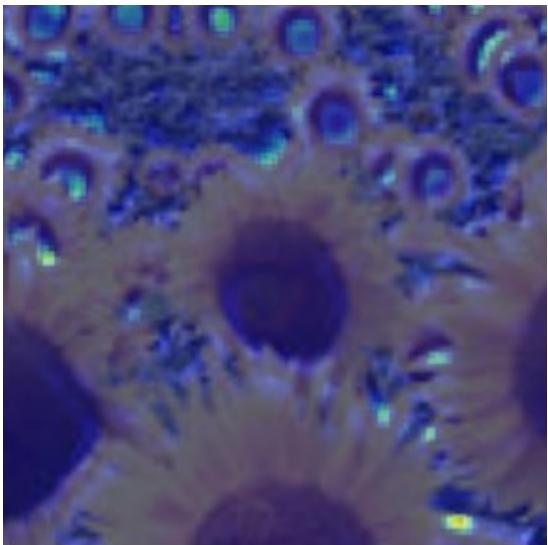
15.5



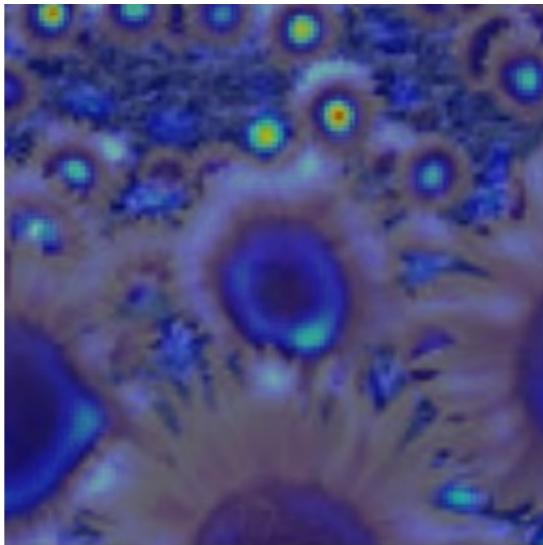
17.0



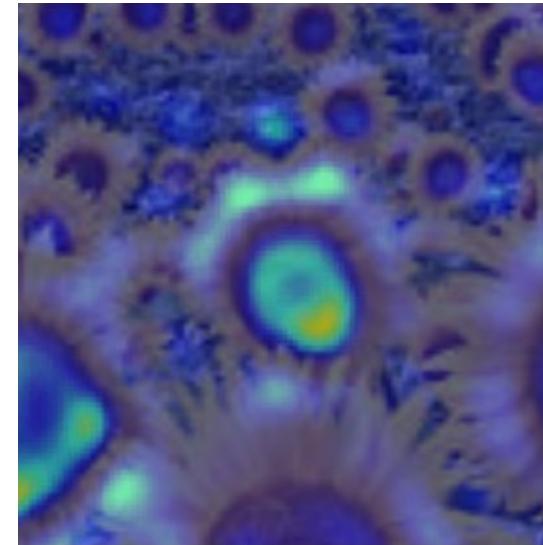
2.1



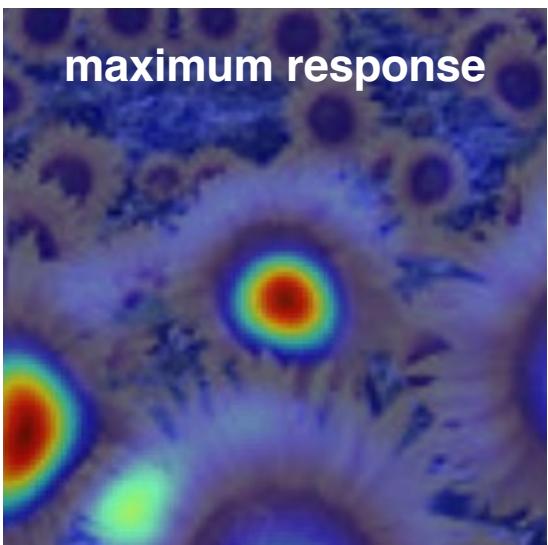
4.2



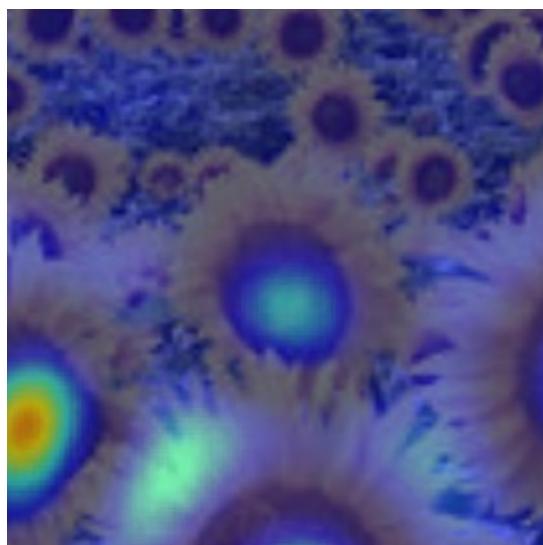
6.0



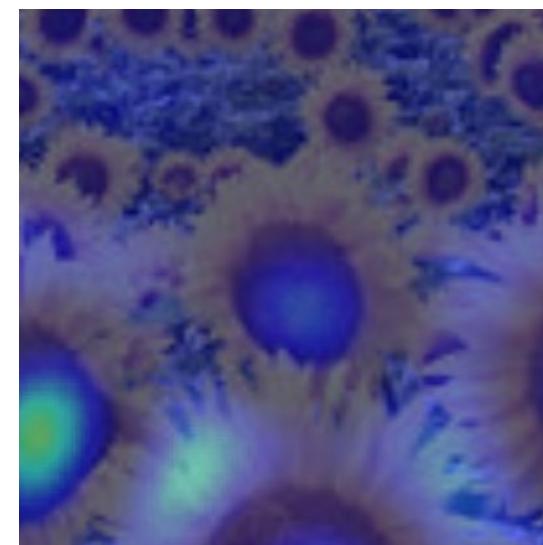
9.8



15.5

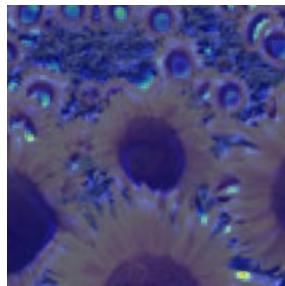


17.0

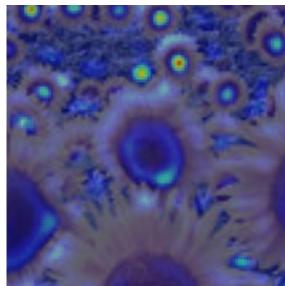


# optimal scale

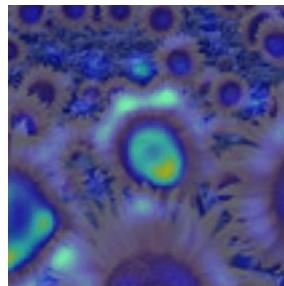
2.1



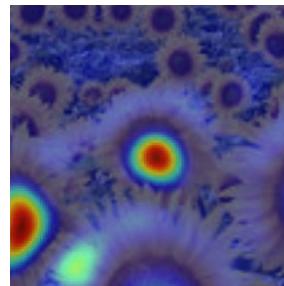
4.2



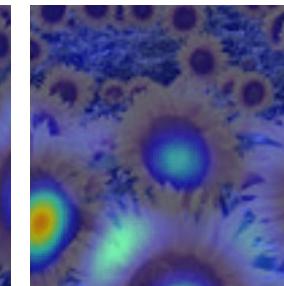
6.0



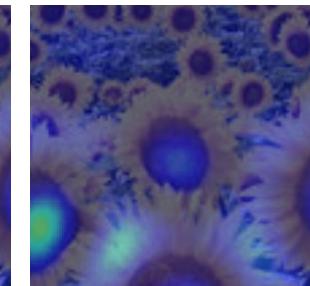
9.8



15.5

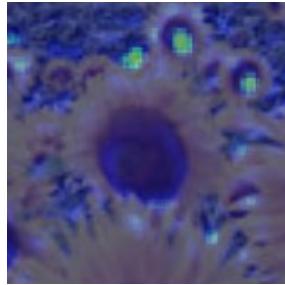


17.0

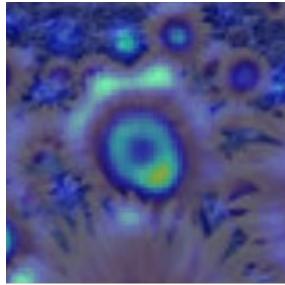


Full size image

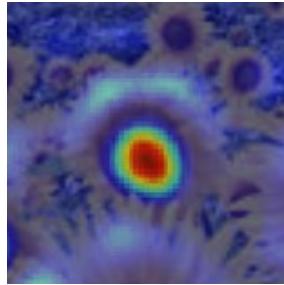
2.1



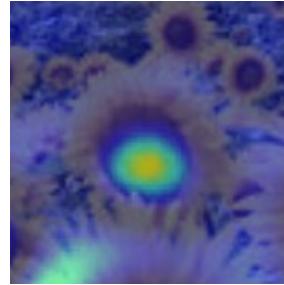
4.2



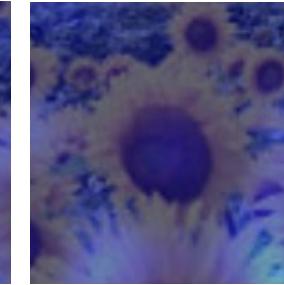
6.0



9.8



15.5

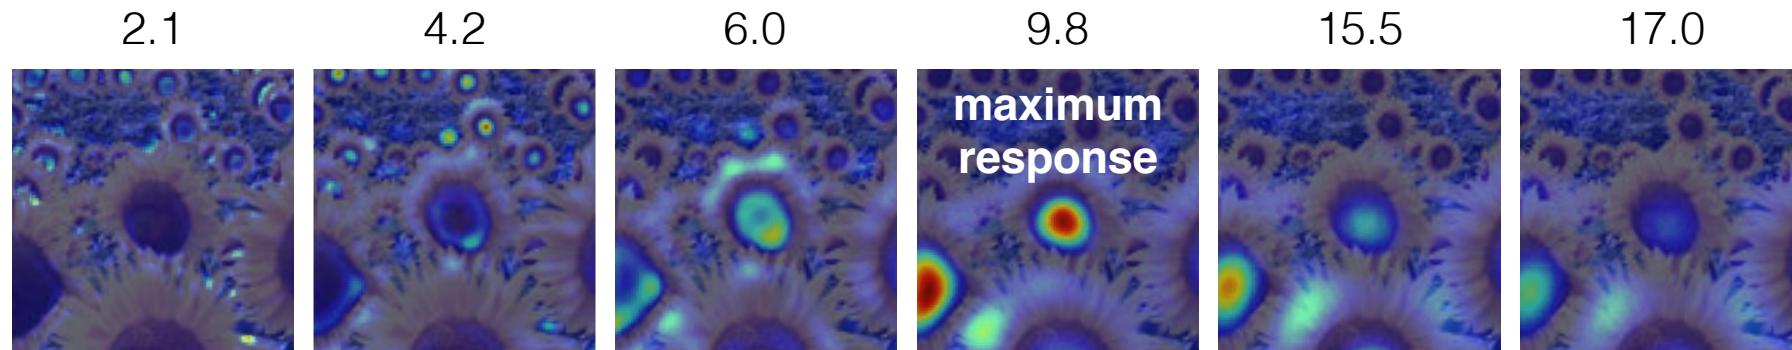


17.0

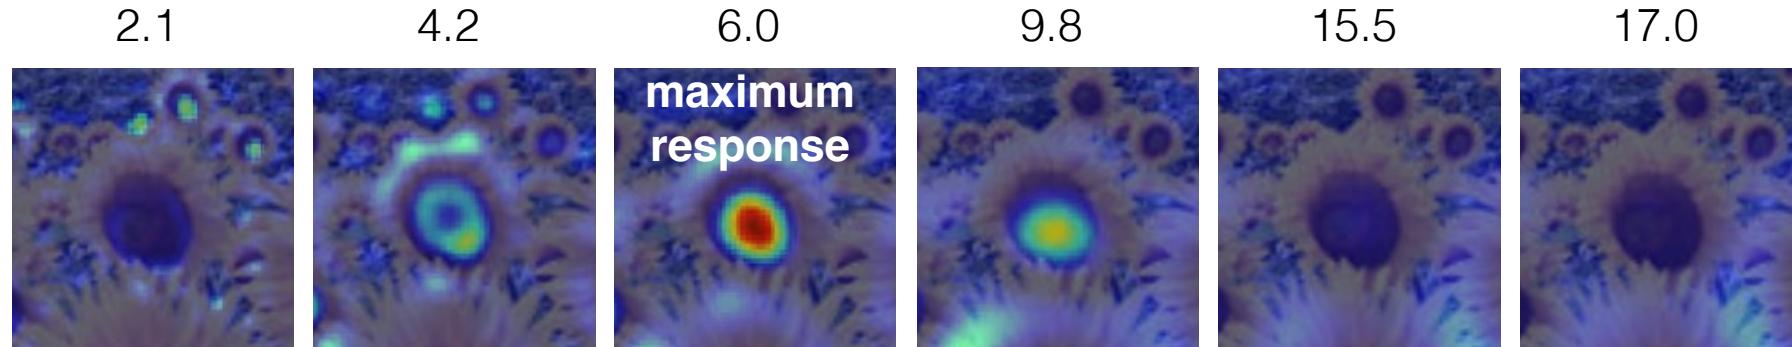


3/4 size image

# optimal scale

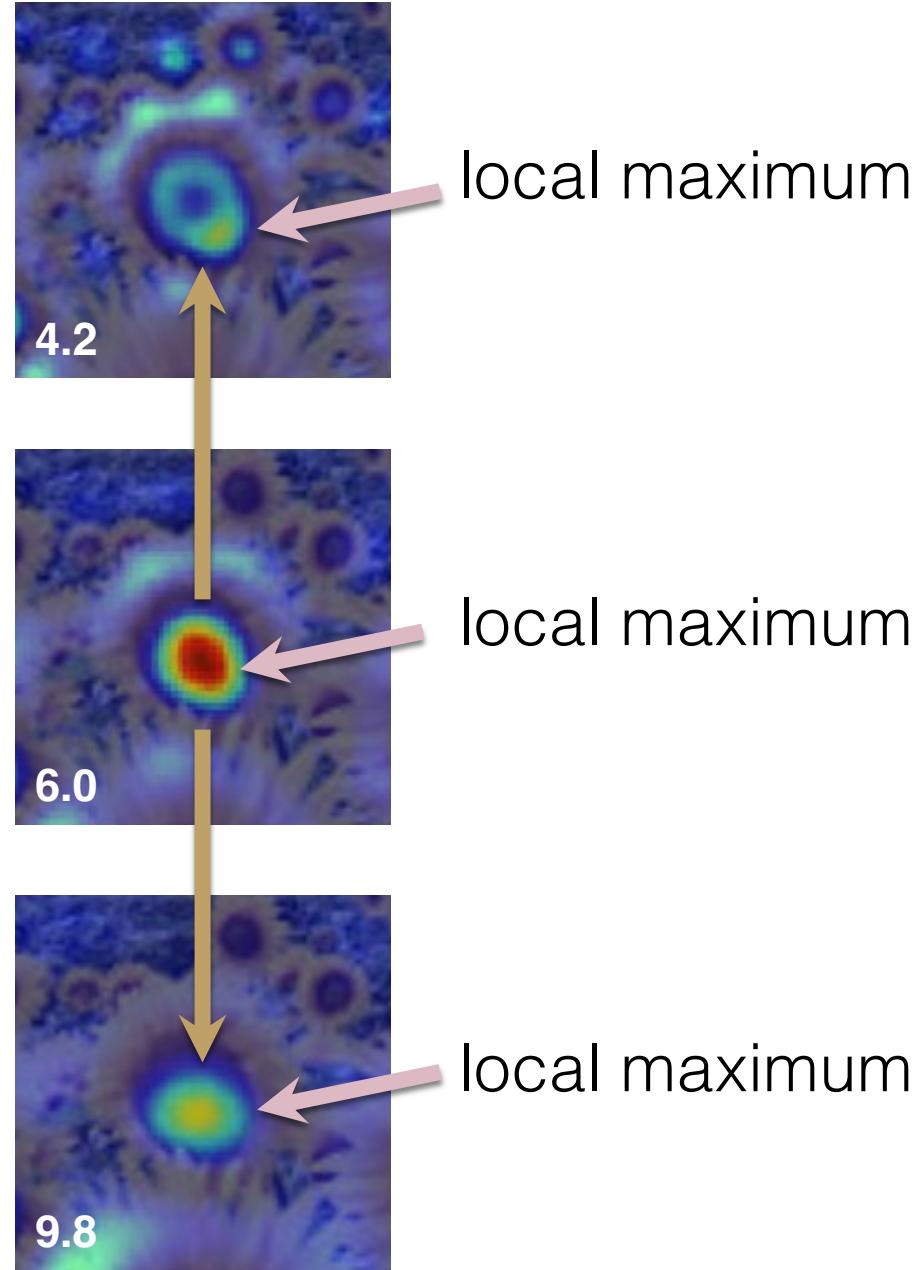


Full size image



3/4 size image

cross-scale maximum



How would you implement scale selection?

# implementation

For each level of the Gaussian pyramid

compute feature response (e.g. Harris, Laplacian)

For each level of the Gaussian pyramid

if local maximum and cross-scale

**save** scale and location of feature  $(x, y, s)$

