


Computer Vision (CS 419/619) Spring 2024

Autoencoders


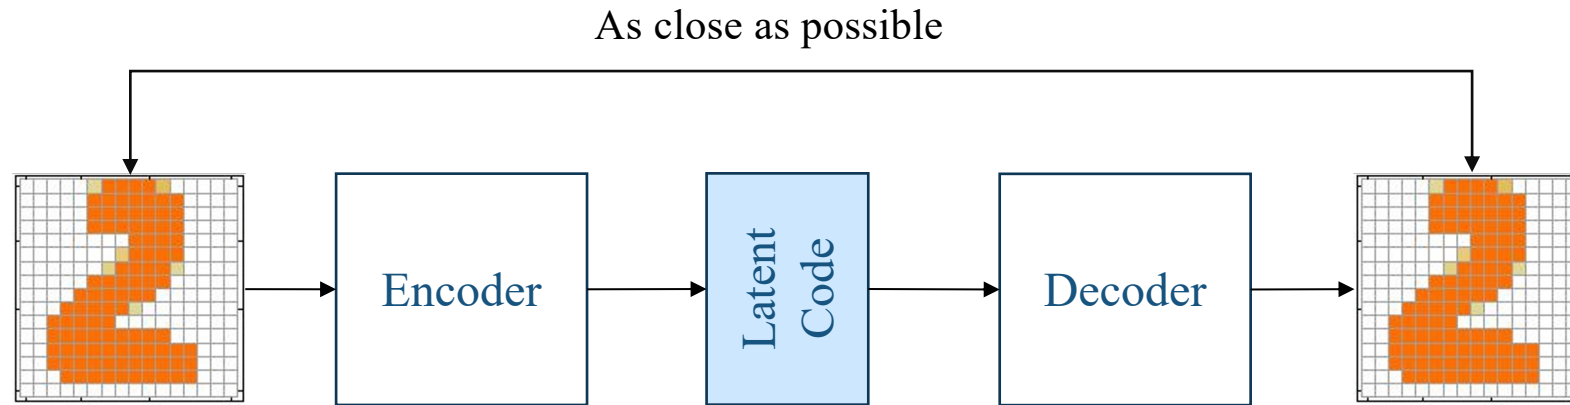


Myth of the Cave (*Plato's Allegory of the Cave*) [1]

- The aim of **generative models** is to learn or understand these latent variables, even when we are just given the observed variables/data.

- The story begins with prisoners who have lived their entire lives chained inside a cave.
- Behind the prisoners is a fire, and between the fire and the prisoners are people carrying puppets or other objects.
- These cast shadows on the opposite wall.
- The prisoners watch these shadows, believing this to be their reality as they have known nothing else.
- From the prisoners' perspective, the shadows are the "observed variables".
- These variables can be perceived and measured.
- Behind them are the true objects that are casting the shadow.
- From the prisoners' perspective, these objects are hidden from them.
- They are can termed as "latent variables".


Generating new images





Generating new images

Notice the features along each axis:



- Along the x-axis if move from 0 to 600.
 - There is a shift in feature
 - The images tend to go from circle to a straight line
- Along the y-axis
 - The above values tend (~ 600) to show a loop on the top of the digit
 - The above values (~ 0) tend to show a loop on the bottom of the digit.
- Grid of values in the latent space.
- The axes denotes the values in the latent vector.
- The images (*coordinates*) denote the reconstructed image from that vector.



Generating new images





Generating new images





Generating new images

- Autoencoder tries to model the probability density of location of each class.
- Some ‘*special*’ types of autoencoder try to model only the parameters of probability distribution for each class.
- *We will revisit the second point in a bit.*











Decoder






Decoder









- The '*special*' Autoencoder tries to learn the location of the '*clusters*' instead of the vectors.
- Learning the clusters:
 - Center of the cluster (μ)
 - Spread of the cluster (Σ)



- The ‘*special*’ Autoencoder tries to learn the location of the ‘*clusters*’ instead of the vectors.
- Learning the clusters:
 - Center of the cluster (μ)
 - Spread of the cluster (Σ)



$$\mu = \begin{bmatrix} 0 \\ 0 \end{bmatrix}, \Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \quad \mu = \begin{bmatrix} 0 \\ 0 \end{bmatrix}, \Sigma = \begin{bmatrix} 1 & 0.5 \\ 0.5 & 1 \end{bmatrix} \quad \mu = \begin{bmatrix} 0 \\ 0 \end{bmatrix}, \Sigma = \begin{bmatrix} 1 & 0.8 \\ 0.8 & 1 \end{bmatrix}$$



1. Wikipedia,
https://en.wikipedia.org/wiki/Gaussian_function

Reference and study material:

1. Deep Learning - Part 1, Dr. Mitesh M. Khapra, NPTEL, <https://archive.nptel.ac.in/noc/courses/noc19/SEM2/noc19-cs85/>
2. Images are taken from Pixabay, <https://pixabay.com/>